

MobileDeepfake: A Lightweight Cascaded Deepfake Detection System for Mobile Deployment

Wing Yam Ho
The Hong Kong Polytechnic University
23041062

November 26, 2025

Abstract

Deepfakes threaten media integrity and platform trust, yet many detectors that perform well on curated benchmarks degrade sharply under distribution shift. For practical use, we target on-device detection where latency, footprint, and privacy constraints are strict, and where missed fakes (false negatives) carry disproportionate risk. We present MobileDeepfake, a two-inference cascade: a lightweight MobileNetV4 makes high-confidence decisions quickly, while uncertain samples are escalated to a higher-accuracy EfficientNetV2 expert. The system is trained across multiple datasets and tuned with a cost-aware threshold grid search that explicitly minimizes false negatives under an escalation-rate budget, and we export compact quantized artifacts for mobile integration.

Our pipeline is reproducible end-to-end: data manifests and loaders enable multi-source training; optional hard-example mining ablations concentrate the Stage 2 expert on ambiguous cases, while the default released training script uses uniform sampling and corresponds to our main reported results; and a tuning/evaluation suite generates the tables and figures used in the paper. On combined validation, the EfficientNetV2 expert achieves high AUC. The tuned cascade attains high AUC while driving FNR to low, exposing an interpretable recall-vs-compute trade-off. We assess robustness under common perturbations (JPEG, noise, blur, brightness) and apply temperature scaling for probability calibration when thresholds transfer across datasets. Cross-dataset evaluation on Deepfake-Eval-2024 highlights the generalization gap typical of in-the-wild media, reinforcing the need for calibration and domain-aware tuning.

We focus on deployability rather than proposing a new backbone: the contribution is a practical, auditable methodology that couples a two-stage cascade with explicit threshold controls, calibration, and a mobile export path. Released artifacts and scripts allow results to be reproduced and extended (e.g., swapping the expert, adding distillation/QAT, or adopting per-family thresholding for robustness). This work provides a template for balancing recall, efficiency, and maintainability in mobile deepfake detection.

1 Introduction

Deepfakes pose growing risks to trust, safety, and platform integrity. Synthetic face-swap videos and images, powered by increasingly accessible generative models, are now routinely weaponized for harassment, disinformation, financial fraud, and non-consensual intimate content. Recent incidents highlight the real-world consequences: politically motivated deepfakes have influenced elections [1], celebrity deepfakes have enabled large-scale scams, and synthetic identity fraud has cost financial institutions millions. Unlike traditional image manipulations, modern deepfakes—generated by

GANs, diffusion models, or advanced reenactment systems—can achieve photorealistic quality that fools both human observers and naive automated detectors.

At the same time, media consumption has shifted dramatically toward mobile devices. Mobile phones account for over 60% of internet traffic globally, with platforms like TikTok, Instagram, and WhatsApp serving billions of users primarily through mobile apps. Consequently, practical deepfake countermeasures must operate robustly and efficiently on resource-constrained phones and tablets, often under limited connectivity and strict privacy requirements. In many deployments—especially in privacy-conscious regions or for offline tools—media never leaves the device and only aggregate detection signals are logged. This on-device constraint eliminates server-side processing as a fallback and requires models to be simultaneously accurate, compact, and fast.

Mobile platforms therefore face a fundamental tension between strong detection accuracy and acceptable user experience. Models must achieve high recall (low false negative rates) to prevent harmful content from evading detection, yet they must also be small enough to fit within app size budgets and fast enough to avoid noticeable latency or excessive battery drain. Under distribution shift—when encountering manipulation methods, compression pipelines, or content types absent from training data—detectors must maintain robustness without escalating to expensive fallback models for every sample. We target this on-device deployment setting and focus on three primary challenges: (i) cross-dataset generalization in the wild, (ii) low latency and small footprint on commodity hardware, and (iii) strict control of false negatives (missed fakes) under operational compute budgets.

From a modeling perspective, recent detectors increasingly adopt Transformer or hybrid CNN-ViT architectures that improve accuracy but raise compute and memory costs, complicating mobile deployment [2, 3]. At the same time, cross-dataset evaluations consistently report notable performance drops from curated benchmarks such as FaceForensics++ and CelebDF-v2 to in-the-wild media (cross-dataset gap) [4, 5, 6, 7]. This gap is further exposed by newer benchmarks such as Deepfake-Eval-2024, which aggregate manipulated content from diverse websites and languages under heterogeneous compression and editing pipelines. For practitioners operating platforms, the central question is not only “What AUC can this model reach on a held-out test set?” but “How often will a malicious fake slip through under realistic load and latency constraints?” In other words, the risk of missed fakes and the compute budget per sample must be explicitly modeled.

Several strands of prior work are relevant to this goal. Deepfake-specific detectors explore spatial, frequency, and physiological cues [7, 8, 9], while lightweight vision models such as MobileNet and EfficientNet families target efficient inference on ARM-class hardware [10, 11, 12, 13]. Separately, early-exit networks and cascaded systems [14, 15, 16] show how to trade compute for accuracy by routing easy samples to shallow branches. However, existing deepfake work rarely combines these ingredients into an end-to-end, auditable system that is (i) trained across multiple datasets, (ii) tuned with explicit risk and cost metrics, and (iii) packaged as a reproducible mobile deployment template. Most studies either evaluate a single heavy detector on server-class GPUs or treat mobile optimization as a brief post-processing step.

In this work we adopt a system-level perspective and design a cascaded detector tailored for on-device deployment. Our goal is to expose explicit knobs over recall, escalation rate, and compute while keeping the training and evaluation pipeline reproducible from raw manifests to mobile artifacts. Against this backdrop, we ask:

- *How can we design a deepfake detector that provides high recall under distribution shift while remaining deployable on resource-constrained mobile devices?*
- *Can a cascaded architecture with explicit thresholds offer interpretable trade-offs between compute and risk, without sacrificing reproducibility?*

- *What minimal set of scripts and artifacts is needed so that practitioners can audit, reproduce, and extend the full pipeline?*
- *How should we evaluate and report robustness—for example under compression, noise, and blur—and cross-dataset generalization in a way that directly informs deployment decisions?*

Existing works on mobile vision and early exits provide building blocks (lightweight backbones, quantization, cascades), but rarely integrate them into an end-to-end, auditable system tailored to deepfake detection with multi-dataset training and out-of-distribution (OOD) evaluation.

Contributions. We design **MobileDeepfake**, a cascaded detection system engineered for mobile deployment and rigorous evaluation. Our contributions are:

- **A cost-aware, two-stage cascade with explicit risk controls.** A lightweight MobileNetV4 Stage 1 acts as a fast filter; ambiguous samples are escalated to an EfficientNetV2 expert (Stage 2), with an optional Stage 3 LightGBM meta-model. We define Stage 1 leakage (the fraction of fakes incorrectly accepted at Stage 1) and Stage 2 escalation rate as primary system metrics, and use a simple two-threshold routing scheme coupled with a cost-sensitive grid search (formalised in Section 5.1) to expose interpretable trade-offs between false negative rate (FNR) and compute.
- **Multi-dataset training and cross-dataset evaluation, including Deepfake-Eval-2024.** We train on four academic datasets and hold out Deepfake-Eval-2024 for OOD validation and test, quantifying performance gaps and highlighting practical failure modes under in-the-wild distribution shift. Our experiments report both single-model baselines and cascaded results, with calibrated vs. uncalibrated comparisons when transferring thresholds across datasets.
- **An end-to-end, scriptable pipeline with hard-example mining, calibration, and robustness analysis.** The system includes data preprocessing, manifest generation, hard-sample mining, expert training, meta-model fitting, and threshold tuning, with scripts that generate the tables and figures used in this paper. Robustness sweeps over common corruptions (JPEG compression, noise, blur, brightness) and a Deepfake-Eval-2024 evaluation stage are integrated into the same pipeline.
- **A mobile-oriented export and integration template.** We provide a mobile optimization stage that applies quantization and knowledge distillation to shrink models, and a deployment stage that exports compact TorchScript artifacts (and ONNX bundles when needed) together with thresholds/temperature in a metadata file suitable for Android integration. While our focus is on methodology rather than a new backbone, the implementation illustrates how to turn deepfake research pipelines into deployable, auditable systems.

Concretely, MobileDeepfake follows a train $\rightarrow$ tune $\rightarrow$ deploy workflow across six stages. The core online inference in our final system is a two-stage cascade: Stage 1 trains a lightweight MobileNetV4 filter for fast screening; Stage 2 builds an EfficientNetV2 expert and produces per-sample logits/features. Stage 3 optionally fits a LightGBM meta-model on K-fold meta-datasets to study fusion but is not part of the default deployment. Stage 4 grid-searches routing thresholds to minimize false negatives under an escalation-rate budget and logs Stage 1 leakage alongside Stage 2 usage. Stage 5 conducts cross-dataset and robustness evaluation (including Deepfake-Eval-2024), and Stage 6 exports quantized artifacts for on-device integration.

We focus on deployability rather than proposing a new backbone: the contribution is a practical, auditable methodology that couples a two-stage cascade with explicit threshold controls, calibration,

robustness analysis, and a mobile export path. We do explore stronger hybrid CNN-Transformer experts such as GenConViT in our Stage 2 ensemble, but find that they offer limited and unstable gains relative to EfficientNetV2-B3 under our constraints, so we relegate them to an appendix and base the default cascade on MobileNetV4 + EfficientNetV2. Subsequent sections detail the datasets and task, the cascaded method and optimization objectives, the experimental setup, and empirical results on in-distribution and out-of-distribution data, followed by a discussion of limitations and deployment considerations.

2 Background

Deepfake detection research spans spatial, frequency, temporal, and multimodal cues, as well as dataset curation and evaluation protocols. Recent comprehensive surveys [17, 7, 18] highlight two persistent gaps that motivate our work: (i) detectors that perform well on curated benchmarks often degrade sharply under distribution shift, and (ii) most evaluations assume server-class hardware rather than mobile devices. These gaps underscore the need for system designs that treat cross-dataset generalization and mobile constraints as first-class objectives rather than afterthoughts.

Deepfake Detection as a Problem Space. Deepfake detection is fundamentally a binary classification task: given an image or video frame, determine whether it depicts a real or synthetically manipulated face. The challenge lies in the diversity of manipulation techniques (face swapping, reenactment, attribute editing, full synthesis) and the variability of real-world conditions (compression artifacts, lighting, camera sensors, post-processing). Early approaches relied on handcrafted forensic features such as head pose inconsistencies [19], eye blinking patterns, or JPEG compression artifacts [20]. Modern methods leverage deep learning to extract hierarchical representations directly from pixels or frequency spectra, achieving high accuracy on in-distribution benchmarks [4, 21, 5].

However, cross-dataset generalization remains elusive: models trained on FaceForensics++ or CelebDF can exhibit AUC drops of 20–30% when evaluated on in-the-wild datasets like WildDeepfake [6] or Deepfake-Eval-2024 [1]. This performance gap stems from dataset bias (e.g., training sets overrepresent specific generators or compression settings), domain shift (e.g., social media platforms apply proprietary compression pipelines), and the rapid evolution of synthesis methods. Addressing these gaps requires multi-dataset training protocols that expose models to diverse artifacts, as well as rigorous cross-dataset and out-of-distribution (OOD) evaluation protocols that measure robustness beyond in-distribution accuracy.

Mobile Deployment Constraints. Most deepfake detection research assumes unlimited computational resources—models are evaluated on server GPUs with batch sizes and latencies that are impractical for on-device inference. Yet many real-world applications (e.g., live video verification, client-side content moderation, edge forensics) require deployment on commodity smartphones or embedded devices with constrained memory, power budgets, and thermal envelopes. Mobile deployment introduces additional challenges: (i) model size must fit within typical app bundles (50–100 MB), (ii) per-frame latency must support real-time or near-real-time inference (e.g., <200 ms), and (iii) energy consumption must preserve battery life across extended usage. These constraints motivate the use of lightweight architectures, model compression techniques, and adaptive inference strategies that balance accuracy and efficiency.

Core Detection Approaches. Modern deepfake detectors primarily use convolutional neural networks (CNNs) or vision transformers (ViTs) as feature extractors. CNN-based approaches [4, 22, 11] leverage hierarchical convolutional layers to capture spatial artifacts such as blending boundaries, facial inconsistencies, and generative noise patterns. The MobileNet [23, 24, 25, 10] and EfficientNet [22, 11] families exemplify architectures optimized for mobile and resource-constrained settings through depthwise separable convolutions and compound scaling.

Frequency-domain methods [8, 9, 26] exploit the observation that GAN-generated images often exhibit distinctive patterns in their Fourier or DCT spectra, providing complementary cues that survive spatial-domain compression. Recent work demonstrates that combining spatial and frequency features improves robustness to JPEG compression and other post-processing operations [27, 28].

Vision transformers [2, 3] offer global receptive fields that can capture long-range dependencies and subtle manipulation cues across entire frames. Hybrid CNN-Transformer architectures [29, 30, 31] aim to combine the inductive biases of CNNs (e.g., translation equivariance) with the modeling capacity of transformers, though at the cost of increased computational complexity.

Temporal modeling via 3D CNNs [32, 33], LSTMs [34, 35], or video transformers [36] extends frame-level detection to video sequences, enabling the use of inter-frame consistency and motion artifacts as additional forensic signals. However, temporal approaches typically require multi-frame context windows and are challenging to deploy on mobile devices due to memory and latency constraints.

Model Compression and Efficiency. To bridge the gap between research models and deployable systems, practitioners rely on three primary compression techniques:

- **Pruning** [37, 38, 39]: Removing redundant weights or channels to reduce model size and FLOPs. Structured pruning (removing entire filters) offers better hardware efficiency than unstructured (weight-level) pruning.
- **Quantization** [40, 41, 42]: Representing weights and activations with lower-precision data types (e.g., INT8 instead of FP32). Post-training quantization (PTQ) is fast but limited to 8-bit precision, while quantization-aware training (QAT) enables more aggressive compression (e.g., INT4) at the cost of retraining.
- **Knowledge distillation** [43, 44, 45]: Transferring knowledge from a large teacher model to a compact student model by matching soft probability distributions or intermediate feature representations.

Modern deep learning frameworks (PyTorch, TensorFlow, ONNX Runtime) provide end-to-end toolchains for converting research models into mobile-optimized INT8 variants and exporting them to Android (via ONNX, TFLite) or iOS (via CoreML) runtimes.

Cascade and Adaptive Inference. Cascade systems and early-exit networks [14, 15, 46] offer an alternative efficiency strategy: route easy samples through lightweight classifiers and escalate hard ones to more expensive experts. Classical examples include Viola-Jones face detection [47], which uses a cascade of progressively complex classifiers to reject non-faces early. In deep learning, BranchyNet [14] attaches shallow classifiers to intermediate layers, enabling early termination for confident predictions. Recent work on adaptive computation time (ACT) [48] and dynamic depth networks [49] learns input-dependent routing policies.

However, most cascade work focuses on maximizing accuracy or minimizing FLOPs on fixed test distributions, without exposing interpretable operational metrics (e.g., escalation rate, Stage 1 leakage rate) that platform operators need to tune cost-quality trade-offs. Furthermore, cascade

detectors for deepfake detection remain rare in the literature, and seldom report performance under distribution shift.

Calibration and Distribution Shift. Confidence calibration via temperature scaling [50, 51] aligns predicted probabilities with empirical frequencies, improving reliability diagrams and expected calibration error (ECE). Well-calibrated models are particularly valuable for cascade routing: thresholds set on calibrated probabilities better reflect true uncertainty and generalize more robustly to OOD data. However, calibration alone does not eliminate performance gaps under domain shift [7]; a well-calibrated model with 75% accuracy on one distribution may still exhibit poor calibration if its true accuracy on another distribution is 60%. Addressing cross-dataset gaps fundamentally requires domain adaptation techniques [52, 53, 54], diverse training data, and transparent reporting of OOD performance.

Our Approach. Given these challenges, we propose a two-stage cascade system for mobile deepfake detection that explicitly optimizes for both detection quality and deployment efficiency. Our Stage 1 filter uses MobileNetV4 [10]—a lightweight CNN optimized for mobile hardware—to perform fast frame-level screening. Samples with highly confident Stage 1 predictions (very likely real or very likely fake) receive immediate labels, while ambiguous cases escalate to a Stage 2 expert (EfficientNetV2-B3 [11] or optional GenConViT [31]) for high-fidelity classification. We formulate threshold selection as a cost-sensitive optimization problem: minimize false negative rate (FNR) subject to a bounded Stage 2 escalation budget, exposing interpretable trade-offs between recall and compute cost.

We train both stages on a combined multi-dataset manifest (CelebDF-v2, FaceForensics++, DFDC, DeeperForensics-1.0) to promote cross-dataset generalization, experiment with hard example mining to focus Stage 2 capacity on cases where Stage 1 is uncertain, and validate the cascade on the in-the-wild Deepfake-Eval-2024 benchmark [1] to measure OOD performance. In our main results and public training script, we use the simpler standard configuration with uniform sampling and treat hard-example mining as an ablation. For mobile deployment, we apply post-training quantization (INT8) and export calibrated ONNX models with fixed thresholds, achieving sub-200 ms latency on a Xiaomi 13 smartphone (Snapdragon 8 Gen 2).

The remainder of this paper proceeds as follows: Section 3 surveys related work in detail, positioning our contributions within the broader landscape of deepfake detection, cascade architectures, and mobile deployment. Section 4 describes our datasets and task formulation. Section 5 presents our six-stage methodology, including formal problem definitions, cascade optimization, and hard example mining. Sections 6 and 7 detail experimental setup and results on in-distribution validation and OOD benchmarks. Section 8 reports ablation studies, and Section 9 discusses mobile deployment specifics (ONNX export, on-device latency, accuracy). Finally, Section 10 discusses limitations and future work, and Section 11 concludes.

3 Related Work

Detection Architectures

CNN-Based Approaches

On early deepfake benchmarks such as FaceForensics++ and DFDC, Xception-based CNNs became de facto baselines [4, 21, 5]. EfficientNet [22] and its successor EfficientNetV2 [11] achieved strong performance through compound scaling that optimizes depth, width, and resolution jointly. Recent

studies [34, 55] demonstrate that even lightweight architectures like MobileNetV3 [25] can achieve competitive detection accuracy when properly fine-tuned. ResNet-family models [56], including ResNeXt variants, have been widely adopted for spatial feature extraction in deepfake detection [?]. Informed by these CNN architectures and mobile efficiency considerations, we adopt MobileNetV4 for our Stage 1 filter and EfficientNetV2 for our Stage 2 expert, balancing detection accuracy with on-device latency constraints.

Frequency-Domain Methods

Frequency analysis has emerged as a powerful approach for detecting deepfake artifacts that survive spatial-domain compression. Durall et al. [8] demonstrated that GAN-generated images exhibit distinctive frequency patterns in their spectral distributions. Qian et al. [9] proposed adaptive frequency filtering to emphasize manipulation cues in the frequency domain. More recent work [26, 28] leverages Discrete Cosine Transform (DCT) and Fast Fourier Transform (FFT) to extract high-frequency artifacts as forensic fingerprints. FreqNet [26], introduced at AAAI 2024, applies convolutional layers to both phase and amplitude spectra between FFT and inverse FFT operations, achieving state-of-the-art generalization (+9.8% improvement) across 17 GANs. Studies using Discrete Fourier Transform [57] report classification accuracies exceeding 99% when frequency-domain features are combined with traditional ML classifiers. Wang et al. [27] proposed Adaptive Frequency learning (AFD) that dynamically adjusts frequency emphasis during training to improve robustness under compression. While we focus on spatial-domain CNN architectures for simplicity and inference speed, frequency-domain extensions remain compatible with our cascade framework and could provide complementary robustness under post-processing.

Vision Transformers and Hybrid Architectures

Vision Transformers (ViTs) have recently been adapted to deepfake detection with promising results. Vanilla ViT [2] and hierarchical variants such as Swin Transformer [3] provide global receptive fields that can capture subtle manipulation artifacts across the entire image. A comprehensive 2024 survey [58] categorizes 14 ViT-based deepfake detection models into standalone, sequential (e.g., MARLIN, ISTVT), and parallel (e.g., M2TR) architectures. UIA-ViT [59] achieved 99.33% AUC on FaceForensics++ by incorporating uncertainty-informed attention mechanisms.

Hybrid CNN-Transformer architectures combine the inductive biases of CNNs with the global modeling capacity of transformers. Coccomini et al. [29] combined EfficientNet and vision transformers for video deepfake detection in 2022, while Khan and Dang-Nguyen [30] proposed a Hybrid Transformer Network that fuses convolutional features with transformer encodings. GenConViT [31], which we explore in our Stage 2 expert ensemble, represents a generative convolutional vision transformer that operates in dual modes (hybrid and pretrained). Under our multi-dataset training and mobile deployment constraints, GenConViT matched but did not consistently surpass EfficientNetV2-B3 and exhibited higher training instability, so we treat it as an exploratory component summarised in the appendix rather than part of the default deployed cascade. CoAtNet-based ensembles [60] leverage self-attention mechanisms in bagging frameworks for manipulated face detection. Das et al. [61] introduced masked autoencoding with spatiotemporal transformers for enhanced video forgery detection.

Attention Mechanisms and Multi-Scale Analysis

Attention mechanisms have proven effective in focusing on manipulation artifacts. Multi-attention networks [62] combining region-independent loss and attention-guided data augmentation achieved

state-of-the-art performance across multiple datasets. ReLAF-Net [63] employs restricted self-attention on deep CNN features to learn both fine-grained local relationships and global dependencies. MobileNetV3 augmented with self-attention layers [64] demonstrated that lightweight architectures can benefit significantly from attention mechanisms while maintaining deployment efficiency. Spatiotemporal attention mechanisms [65, 35] have been integrated with ConvLSTM to capture and enhance frame-to-frame correlations in video sequences.

Temporal and Video-Level Detection

Beyond spatial architectures, temporal modeling helps handle partial manipulations and contextual cues in video [6, 66]. 3D CNNs [32] extend 2D convolutional filters to spatiotemporal dimensions, enabling direct learning from video clips. The I3D (Inflated 3D) model [33] modifies 2D Inception networks by inflating convolutional and pooling kernels to 3D. Hybrid CNN-LSTM architectures [34, 67] use CNNs for spatial feature extraction at the frame level, then feed sequential features into LSTMs for temporal modeling. Xception-LSTM with ConvLSTM [35] incorporates spatiotemporal attention before dimension reduction to preserve frame structure information. EfficientNet-Bi-LSTM [68] and similar approaches extract spatial information via EfficientNet and acquire temporal information through bidirectional LSTMs. TimeSformer [36] and other video transformers provide pure transformer-based alternatives for spatiotemporal reasoning.

Parameter-efficient adapters (e.g., DeepFake-Adapter [69]) further improve transfer to in-the-wild data by tuning small sets of adapter parameters on top of a frozen backbone, reducing computational cost while maintaining generalization. Other lines exploit physiological signals (e.g., eye blinking [19], head motion) and camera fingerprints [20], but typically rely on specialized preprocessing pipelines. In this work we focus on frame-level classification for reproducibility and mobile latency, using generic image backbones (MobileNetV4 and EfficientNetV2) and optional GenConvViT experts. Our pipeline remains compatible with temporal post-processing when budgets allow, but deliberately keeps the core inference path simple enough for on-device deployment.

Lightweight Models and Compression

Efficient Architecture Design

Mobile vision emphasizes efficient backbones tailored to on-device constraints. The MobileNet family [23, 24, 25] pioneered depthwise separable convolutions to reduce computational cost while maintaining accuracy. MobileNetV3 [25] incorporated hardware-aware NAS and squeeze-and-excitation modules to optimize for mobile CPUs and GPUs. The latest MobileNetV4 [10] further scales up mobile CNNs with hybrid architectures and improved training recipes. EfficientNet [22] and its successor EfficientNetV2 [11] deliver strong accuracy under tight FLOPs budgets through compound scaling that jointly optimizes network depth, width, and resolution. ShuffleNet [12, 13] and SqueezeNet [70] provide alternative efficient designs that prioritize inference speed on resource-constrained devices.

Hardware-aware neural architecture search (NAS) [71] and once-for-all training [72] enable per-device specialization from a single supernet, allowing practitioners to extract task-and-device-specific sub-networks without retraining from scratch. Recent real-time deepfake detection systems [73, 55] demonstrate that lightweight models such as MobileNetV3 and binary neural networks can achieve competitive performance when combined with attention mechanisms and proper fine-tuning strategies.

Model Compression Techniques

Pruning and sparsity methods reduce model size by removing redundant weights and connections. Structured pruning removes entire channels or filters, while unstructured pruning targets individual weights [39]. Han et al. [37, 38] demonstrated that deep networks can be compressed by 10-50 $\times$ through iterative pruning, quantization, and Huffman coding. The Lottery Ticket Hypothesis [74] suggests that dense networks contain sparse sub-networks that can achieve comparable accuracy when trained from the right initialization. Recent surveys [39] provide comprehensive comparisons of pruning methods and their impact on inference efficiency.

Knowledge distillation [43] transfers knowledge from a large teacher model to a compact student model by matching soft probability distributions or intermediate feature representations. FitNets [44] extend this idea by matching hint layers at intermediate depths, while Born-Again Networks [75] apply multi-generation distillation where students can outperform teachers. Recent surveys [45] synthesize these trends and highlight that distillation is particularly effective when combined with pruning and quantization. Studies on knowledge distillation for edge deployment [76, 77] demonstrate that distilled models can achieve 70% compression with minimal accuracy degradation, making them suitable for fog/edge devices in applications like automatic disease diagnosis and on-device deepfake detection.

Quantization for Mobile Deployment

Quantization has become a central lever for efficient inference, reducing both model size and computational cost by representing weights and activations with lower-precision data types. Post-training quantization (PTQ) [40, 41] converts trained FP32 models to INT8 or lower bitwidths without retraining, offering a fast and lightweight approach to compression. Quantization-aware training (QAT) [41, 78] simulates quantization during training, enabling lower bitwidths (e.g., INT4) with competitive accuracy but requiring access to training data and greater computational expense. A 2021 white paper [42] provides a comprehensive survey of quantization methods, highlighting that PTQ is sufficient for 8-bit quantization with near-FP32 accuracy, while QAT enables aggressive sub-8-bit compression.

Mixed-precision and per-channel quantization schemes [78, 41] provide competitive accuracy at low bitwidths by allowing different layers or channels to use different quantization scales. Recent work on quantizing large language models [79, 80] and vision transformers demonstrates that quantization remains effective even for modern architectures. For deepfake detection, quantization enables deployment on commodity mobile devices with minimal accuracy loss, and modern deep learning frameworks provide end-to-end toolchains for converting research models into compact INT8 variants and exporting them to mobile platforms.

Our Stage 1 uses MobileNetV4 as a fast filter, and Stage 2 adopts EfficientNetV2 and optional GenConViT experts [11, 31]. Stage 4 builds on QAT + knowledge distillation in a teacher–student setup and compares FP32 vs. INT8 models in terms of accuracy, size, and latency, while Stage 6 exports TorchScript and ONNX artifacts suitable for mobile integration. We emphasize measurable footprint and latency on device: model size, on-device escalation rate, and duty cycle jointly determine energy and user experience, so we report escalation alongside accuracy when comparing options.

Cascade and Adaptive Inference

Early-Exit Networks and Dynamic Depth

Early-exit and cascaded systems trade accuracy for efficiency by routing easy samples to fast exits and escalating hard ones. BranchyNet [14] pioneered this approach by attaching shallow classifiers to intermediate layers of deep networks, enabling early termination for confident predictions. Multi-Scale Dense Networks (MSDNet) [15] extend this idea by learning feature representations at multiple scales, allowing classifiers at different depths to leverage complementary information. Shallow-Deep Networks [16] address network overthinking by identifying when deeper layers degrade performance on easy samples.

A comprehensive 2024 survey on early-exit deep neural networks [46] categorizes methods into static (BranchyNet, MSDNet) and adaptive (ACT, SkipNet) approaches. Adaptive Computation Time (ACT) [48] for RNNs and SkipNet [49] for CNNs learn dynamic routing policies that adapt network depth to input complexity. Recent work on window-based early-exit cascades [81] incorporates uncertainty estimation for improved reliability. QuickNets [82] introduced a novel cascaded training algorithm where each successive layer is trained only on samples misclassified by previous layers, preventing overconfidence and improving calibration.

Routing Policies and Threshold Selection

Confidence-, entropy-, and margin-based routing are common strategies for determining when to exit early. Hendrycks and Gimpel [83] established maximum softmax probability as a simple baseline for detecting misclassified examples and out-of-distribution samples. Cost-sensitive learning frameworks [84] formalize the trade-off between classification errors and computational costs, enabling threshold selection that minimizes expected cost rather than maximizing accuracy alone. Recent work increasingly frames threshold selection as a multi-objective optimization problem that balances error rates (false positives, false negatives) against compute budgets (FLOPs, latency, energy) [85, 86].

Cascaded systems for object detection [47] and classification have been deployed successfully in production environments, but typically employ hand-tuned thresholds rather than learned policies. Bayesian optimization [87] and grid search remain popular for threshold tuning when the search space is small. For multi-stage systems, joint optimization of all thresholds is often intractable, motivating greedy stage-wise tuning or heuristic policies.

Cascades in Deepfake Detection

In deepfake detection, cascades are less explored than single-model detectors, and are rarely paired with explicit system-level metrics such as Stage 1 leakage rate and Stage 2 escalation rate. Most prior work reports AUC/F1 on a fixed test set without exposing the operational trade-offs platform operators care about. Classical cascade detectors like Viola-Jones [47] for face detection demonstrated the effectiveness of multi-stage architectures, but deepfake detection has largely focused on maximizing single-model accuracy rather than optimizing inference efficiency.

Our system follows the early-exit lineage but makes the design explicitly cost-aware: a lightweight Stage 1 filter (MobileNetV4) sits in front of a stronger Stage 2 expert (EfficientNetV2 + GenConViT) and optional Stage 3 meta-model (LightGBM), and a two-threshold routing scheme defines a tunable “uncertainty band” for escalation. Stage 4 performs grid search over low/high thresholds using cached logits, optimizing for low false negative rate under an escalation-rate budget (typically 1–5%). The resulting operating points are reported with both FNR and Stage 2

escalation rate, exposing interpretable trade-offs between risk and compute. For deployment, we keep the runtime policy intentionally simple: thresholds are tuned offline via grid search and then fixed on device for stability, avoiding the complexity of learned routing policies that may drift under distribution shift.

Generalization and Robustness

Cross-Dataset Generalization

Cross-dataset generalization is a core challenge in deepfake detection: large AUC drops (often 20–30%) are observed when models trained on curated datasets like FaceForensics++ are evaluated on in-the-wild media [6, 7]. State-of-the-art models achieve 96–99% AUC on in-dataset tests but degrade to 72–78% on cross-dataset evaluations [?, 88]. This performance gap stems from dataset bias, compression artifacts, and differences in manipulation techniques across generators.

Domain Adaptive Batch Normalization (DABN) [53] mitigates distribution gaps by utilizing test-set statistics to restore batch normalization layer effectiveness, achieving nearly 20% accuracy improvement on CelebDF under black-box settings. Recent work on unsupervised domain adaptation [54] transfers deepfake detection knowledge from labeled source domains to unlabeled target domains using Domain Distance Optimization (DDO) modules that align inter-domain and intra-domain feature distances. Spatial-frequency collaborative learning methods [89] outperform frequency-aware baselines by 9.5% AUC on CelebDF through joint spatial and frequency domain modeling.

Domain adaptation surveys [90, 52] provide comprehensive taxonomies of transfer learning and domain-adversarial approaches. Domain-Adversarial Neural Networks (DANN) [52] learn domain-invariant features through adversarial training between feature extractors and domain classifiers. However, achieving consistent cross-dataset performance remains elusive: fairness-aware methods [91] maintain fairness metrics on some test sets but fail on others, highlighting the difficulty of generalizing both accuracy and fairness simultaneously. In our work, we train on a combined multi-dataset manifest and evaluate on Deepfake-Eval-2024 as an explicit OOD benchmark to directly measure this generalization gap and report cross-dataset performance honestly.

Robustness to Perturbations and Corruptions

Robustness to common corruptions (JPEG compression, noise, blur, illumination changes) is likewise uneven; even strong detectors can degrade substantially under realistic post-processing [8, 9, 27]. Frequency-domain methods provide one avenue for robustness by emphasizing manipulation cues that survive compression and noise. Adaptive frequency learning [27] dynamically adjusts frequency emphasis during training to improve robustness under JPEG compression. High-frequency enhancement networks [92] for heavily compressed content demonstrate that explicit frequency modeling can recover detection performance even under aggressive compression.

Adversarial robustness is another critical dimension. Carlini and Wagner [93] attacks and Projected Gradient Descent (PGD) [94] perturbations can fool deepfake detectors with imperceptible modifications. While adversarial training [94, 95] improves robustness to ℓ_p -bounded attacks, it often degrades performance on clean samples and does not guarantee robustness to adaptive attacks. Recent comprehensive evaluations [96] assess 8 state-of-the-art defenses against evolving threats (adversarial examples, OOD samples, novel generators), finding that no single method achieves robustness across all threat models.

Benchmarks and Evaluation Protocols

Newer benchmarks explicitly target generalization gaps. Deepfake-Eval-2024 [1] aggregates 1,072 real and 964 fake files from diverse platforms, languages, and compression pipelines as an in-the-wild benchmark. DeepfakeBench [88], presented at NeurIPS 2023, provides a comprehensive evaluation framework supporting 36 detection methods (28 image + 8 video detectors) and introduces the DF40 dataset comprising 40 distinct deepfake techniques including state-of-the-art 2024 generators. Celeb-DF++ [97] extends the original CelebDF benchmark to evaluate 19 classic methods and 5 recent (post-2024) detectors.

Recent surveys [7, 18, 98] emphasize that robustness techniques must be coupled with rigorous cross-dataset evaluation protocols to avoid overfitting to specific artifacts. The surveys categorize methods by feature extraction (spatial, frequency, temporal, multimodal) and network architecture (CNN, Transformer, hybrid), and identify generalization as the primary open challenge for practical deployment.

Calibration and Confidence Estimation

Calibration and thresholding play an important role under distribution shift. Probability calibration techniques such as temperature scaling [50, 51] help align predicted probabilities with empirical frequencies, improving reliability diagrams and expected calibration error (ECE). However, calibration does not eliminate performance gaps on new datasets [7]; a well-calibrated model with 75% accuracy remains poorly calibrated if evaluated on a distribution where its true accuracy is 60%.

Ensemble and meta-learning approaches [99, 100] leverage stacking-based fusion to improve both accuracy and calibration. Naskar et al. [99] combined Xception and EfficientNetV2-B7 features via a meta-learner (MLP), achieving 93% accuracy on CelebDF and 98% on FaceForensics++ through feature ranking and selective fusion. Our Stage 3 meta-model follows this stacking paradigm but uses LightGBM as the meta-learner and incorporates K-fold cross-validation for robustness.

In our evaluation (Stage 5), we therefore combine cross-dataset experiments on Deepfake-Eval-2024 with robustness sweeps over JPEG, noise, motion blur, brightness, downscaling, and gamma perturbations. We report both calibrated cascade results and a Stage 2-only upper bound, and we log Stage 2 escalation rates alongside F1/FNR. This design emphasizes system-level trade-offs: robustness gains are only useful if they do not explode compute or latency on device, and calibration is only useful when interpreted jointly with escalation and leakage statistics rather than as a standalone score.

Deployment and Tooling

Production deployment for mobile requires artifact export, footprint accounting, latency measurement, and minimal telemetry. TorchScript- or ONNX-style runtimes provide stable toolchains for shipping models to Android and iOS, while quantization tooling and surveys [42] make it feasible to convert research models into compact INT8 variants without rewriting training code. Beyond the model itself, practitioners need configuration management (e.g., thresholds, calibration temperatures), logging and monitoring hooks, and a clear separation between offline tuning and online inference.

Our pipeline integrates these considerations end-to-end. We export TorchScript artifacts for Stage 1 and Stage 2 and a bundle metadata file that records thresholds, calibration temperatures, and version information. ONNX exporters and benchmarking scripts provide optional support for alternative runtimes and detailed latency/memory measurements. We also minimise I/O and allocations on device to reduce jitter and simplify maintenance. This positions **MobileDeepfake**

as a bridge between research prototypes and production systems: we emphasise (i) multi-dataset training and hard-example mining; (ii) cost-sensitive threshold tuning for low false negatives with controlled escalation; and (iii) on-device integration with measurable trade-offs. Unlike single-model detectors, our system exposes designer-visible dials (thresholds, temperature) aligned with operational requirements and supported by reproducible scripts and artifacts.

4 Datasets and Task

Task Definition. We frame deepfake detection as binary classification at the frame/image level: *real* (label=0) vs. *fake* (label=1). For videos, our manifests reference extracted face crops; training and evaluation operate at the frame/image granularity to maximise sample size and label fidelity, while downstream applications may aggregate predictions per video (e.g., majority vote or max probability) if needed. This choice keeps the core task simple and aligns with common evaluation practice on FF++/CelebDF/DFDC [4, 5, 21].

Table 1: Dataset sizes by split (rows counted from manifests).

Split	Dataset	Rows
Train	CelebDF-v2	83,599
	DeeperForensics	844,396
	DFDC	721,946
	FaceForensics++	223,653
Val	CelebDF-v2	17,478
	DeeperForensics	165,854
	DFDC	154,919
	FaceForensics++	50,945
Test	CelebDF-v2	18,037
	DeeperForensics	172,894
	DFDC	154,511
	FaceForensics++	47,430
Val	Deepfake-Eval-2024	451,917
Test	Deepfake-Eval-2024	402,413

Table 2: Per-dataset class balance by split (rows and real/fake counts from balanced manifests).

Split	Dataset	Rows	Real	Fake
Train	CelebDF-v2	83,599	41,674	41,925
	DeeperForensics	844,396	414,512	429,884
	DFDC	721,946	361,229	360,717
	FaceForensics++	223,653	111,161	112,492
Val	CelebDF-v2	17,478	8,765	8,713
	DeeperForensics	165,854	87,807	78,047
	DFDC	154,919	77,350	77,569
	FaceForensics++	50,945	26,351	24,594
Test	CelebDF-v2	18,037	9,118	8,919
	DeeperForensics	172,894	89,253	83,641
	DFDC	154,511	77,109	77,402
	FaceForensics++	47,430	23,502	23,928
Val	Deepfake-Eval-2024	451,917	302,369	149,548
Test	Deepfake-Eval-2024	402,413	216,288	186,125

Datasets. We train on four academic datasets—CelebDF v2, FaceForensics++, DeeperForensics 1.0, and DFDC—and hold out Deepfake-Eval-2024 as an out-of-distribution (OOD) benchmark. CelebDF v2 provides high-quality face-swap videos of celebrities designed to challenge detectors [5]; FaceForensics++ aggregates multiple manipulation methods (DeepFakes, Face2Face, FaceSwap, NeuralTextures, etc.) under controlled conditions [4]; DFDC is a large-scale crowd-sourced dataset with diverse actors and compression pipelines [21]; DeeperForensics 1.0 emphasises real-world perturbations and challenging textures [101]. Deepfake-Eval-2024 adds contemporary in-the-wild content collected from TrueMedia.org and social media platforms in 2024, covering 88 websites and 52 languages [1]. We use Deepfake-Eval-2024 strictly for cross-dataset validation/testing. We additionally inspect the DF40 dataset introduced in DeepfakeBench [88] only to determine a target 256×256 PNG specification for our preprocessing pipeline; DF40 itself is not included in the training or evaluation splits for the experiments reported in this paper.

Preprocessing Pipeline. All video datasets are preprocessed into face-crop images using a configurable pipeline (cf. [4, 5, 40] for similar face-crop protocols). Frames are sampled at a fixed interval (`frame_interval=10`), faces are detected with MTCNN (minimum face size `min_face_size=20`, thresholds `[0.6, 0.7, 0.7]`), cropped around the bounding box with `bbox_scale=1.3`, and resized to 256×256 pixels. We cap redundancy by keeping at most 50 faces per video (`max_faces_per_video=50`). These parameters are versioned alongside the manifests so that the training set can be regenerated or audited.

The preprocessing stack supports multiple video backends (e.g., OpenCV, TorchVision, Decord) and face detectors (MTCNN, InsightFace, YOLOv8, MediaPipe, RetinaFace), but our final experiments use MTCNN with GPU-accelerated decoding for simplicity and reproducibility. Processed images are organised into train/validation/test splits for each dataset, with separate folders for real and fake faces, and accompanied by CSV manifests that record the necessary metadata for every crop. This layout keeps the underlying file structure simple while making manifests the primary interface for training, evaluation, and audit.

Manifest Generation and Structure. After preprocessing, we generate CSV manifests that serve as the canonical index for all experiments. Each manifest row corresponds to a single face

image and includes at least `image_path`, `label`, `split`, and `dataset`; additional fields such as `md5`, `valid`, `width`, `height`, and `file_size` facilitate integrity checks. Table 1 is computed directly from these manifests. All sampling, hard-example mining, and evaluation components operate on manifest rows rather than raw folders, making it easy to extend the pipeline to new datasets by regenerating manifests.

Split Protocol and Balancing. We adopt a simple and reproducible split protocol. For each dataset, preprocessed faces are partitioned into train/val/test splits according to configuration ratios, preserving official test lists when available (e.g., CelebDF v2 and FF++). We then construct balanced manifests per dataset and split by sampling equal numbers of real/fake faces. For multi-dataset training, we concatenate these manifests and treat each dataset as a source domain; the combined manifest is approximately balanced both by class and by dataset, so Stage 1 and Stage 2 loaders can rely on standard shuffling rather than complex per-batch reweighting. An earlier weighted sampler remains available in the codebase but is not used in our final experiments.

To guard against trivial path-based leakage, the manifest pipeline can optionally sanitise or reject samples with tell-tale substrings (e.g., directory names that encode `/fake/` or `/real/`). These checks are primarily useful when integrating new corpora; for the curated datasets considered here, manifests are derived from neutral directory layouts and the leakage check is disabled by default.

Augmentation and Normalization. Stage 1 employs a compact torchvision transform pipeline that resizes images to 256 px, applies random horizontal flip, colour jitter, small random affine transformations, and a light Gaussian blur, and normalises with ImageNet statistics. Augmentation is symmetric across classes, with a modest probability of colour/geometry perturbations. This design follows contemporary data augmentation practice for image classification while avoiding heavy corruptions that could distort robustness analysis [102, 7]. Stage 2 uses slightly stronger geometric and colour jitter when beneficial; all stages share the same input resolution and normalisation to simplify deployment and quantization.

OOD Evaluation Protocol. For cross-dataset evaluation, we follow Stage 5 and treat Deepfake-Eval-2024 as a held-out OOD benchmark [1]. We evaluate (i) a calibrated two-stage cascade (Stage 1 + Stage 2 + thresholds tuned on the in-distribution validation set) and (ii) a Stage 2-only upper bound that bypasses Stage 1 and uses the expert on all samples. The tables (Table 8, 9) summarise accuracy, F1, FNR, FPR, and Stage 2 rate on the Deepfake-Eval-2024 val/test splits. The sizable performance gap relative to in-distribution validation underscores the importance of robustness studies and future domain adaptation; in this work we focus on low-friction mitigations such as temperature scaling and threshold sweeps [50, 51], and avoid OOD-specific fine-tuning unless justified by target data.

Dataset Characteristics. Table 3 summarises qualitative properties of each dataset to motivate our protocol design. These traits justify balanced sampling across sources and the use of a dedicated OOD set to probe generalization under realistic distribution shift.

Mobile Evaluation Subset. To evaluate on-device performance without requiring the full multi-gigabyte datasets on a smartphone, we construct a compact `mobile_eval` subset. This subset is derived by stratified sampling from the four training datasets (CelebDF-v2, FaceForensics++, DFDC, DeeperForensics 1.0), selecting representative face crops from each source. The sampling

Table 3: Qualitative characteristics of datasets (informal summary).

Dataset	Generation	Quality	Source	Notes
CelebDF v2	Face-swap	High	YouTube	Celebrities; curated videos
FaceForensics++	Multiple	Mid	Actors	Multiple manipulation families
DFDC	Multiple	Mid-Low	Internet	Large-scale, mixed quality
DeeperForensics 1.0	Multiple	Mid	Lab	Realistic perturbations, hard cases
Deepfake-Eval-2024	Multiple	Mid-Low	Web	In-the-wild, multi-source, 2024 media

procedure picks 1–2 frames per video across a diverse set of videos, targeting approximately 150–200 images total.

The resulting subset is organised into two directories: `images_real/` (72 images) and `images_fake/` (80 images), with ground-truth labels determined by directory membership. This structure allows the Android app to infer labels from file paths during folder detection, enabling automated accuracy computation without manual annotation.

The `mobile_eval` subset serves several purposes: (1) it enables rapid iteration during app development without transferring large datasets; (2) it provides a reproducible benchmark for on-device latency and accuracy measurements; (3) its size is small enough to permit multiple complete runs within a single testing session. The subset intentionally reuses preprocessed face crops from the PC pipeline, ensuring that on-device and desktop evaluations operate on pixel-identical inputs (modulo any JPEG re-encoding during transfer). All on-device metrics reported in Section ?? derive from this subset.

5 Method

5.1 Problem Formulation

Binary Classification Task. We frame deepfake detection as binary classification over face images. Given an input image $x \in \mathcal{X} = \mathbb{R}^{H \times W \times 3}$ (where $H = W = 256$ in our pipeline), the task is to predict a label $y \in \{0, 1\}$ where $y = 0$ denotes real and $y = 1$ denotes fake. Our training data comprises $\mathcal{D} = \{(x_i, y_i, d_i)\}_{i=1}^N$, where d_i indexes the source dataset (CelebDF-v2, FaceForensics++, DFDC, or DeeperForensics-1.0).

Two-Stage Cascade Architecture. Our deployed system consists of two sequential classifiers:

- **Stage 1 (Filter):** A lightweight model $f_1 : \mathcal{X} \rightarrow \mathbb{R}$ producing logit $f_1(x)$ and calibrated probability $p_1(x) = \sigma(f_1(x))$, where σ is the sigmoid function.
- **Stage 2 (Expert):** A higher-capacity model $f_2 : \mathcal{X} \rightarrow \mathbb{R}$ producing logit $f_2(x)$ and calibrated probability $p_2(x) = \sigma(f_2(x)/t)$, where t is the temperature scaling parameter.

Cascade Decision Policy. The cascade routes samples based on two thresholds $\tau_{\text{low}}, \tau_{\text{high}} \in [0, 1]$ with $\tau_{\text{low}} < \tau_{\text{high}}$:

$$\hat{y}(x) = \begin{cases} 0 & \text{if } p_1(x) < \tau_{\text{low}} \\ 1 & \text{if } p_1(x) > \tau_{\text{high}} \\ \mathbf{1}[p_2(x) > \tau_2] & \text{otherwise} \end{cases} \quad (1)$$

where $\mathbf{1}[\cdot]$ is the indicator function and $\tau_2 \in [0, 1]$ is the Stage 2 decision threshold. Samples falling in the interval $[\tau_{\text{low}}, \tau_{\text{high}}]$ are *escalated* to Stage 2, while those outside this band receive immediate Stage 1 predictions.

Evaluation Metrics. We define standard classification metrics on a test set $\mathcal{D}_{\text{test}}$:

$$\text{FNR}(\tau) = \frac{\sum_{(x,y) \in \mathcal{D}_{\text{test}}: y=1} \mathbf{1}[\hat{y}(x) = 0]}{\sum_{(x,y) \in \mathcal{D}_{\text{test}}: y=1} 1} \quad (2)$$

$$\text{FPR}(\tau) = \frac{\sum_{(x,y) \in \mathcal{D}_{\text{test}}: y=0} \mathbf{1}[\hat{y}(x) = 1]}{\sum_{(x,y) \in \mathcal{D}_{\text{test}}: y=0} 1} \quad (3)$$

The **Stage 2 escalation rate** measures the fraction of samples requiring expert evaluation:

$$\pi_2(\tau_{\text{low}}, \tau_{\text{high}}) = \Pr(\tau_{\text{low}} \leq p_1(x) \leq \tau_{\text{high}}) \quad (4)$$

This metric directly governs computational cost, as Stage 2 inference is significantly more expensive than Stage 1.

Optimization Objective. Our cascade tuning (detailed in subsequent subsections) seeks threshold configurations that minimize FNR subject to an escalation budget:

$$\min_{\tau_{\text{low}}, \tau_{\text{high}}, \tau_2} \text{FNR}(\tau_{\text{low}}, \tau_{\text{high}}, \tau_2) \quad \text{subject to} \quad \pi_2(\tau_{\text{low}}, \tau_{\text{high}}) \leq B \quad (5)$$

where B is a platform-defined budget (e.g., $B = 0.05$ limits escalation to 5% of samples). This formulation captures the core trade-off between detection quality (low FNR) and inference cost (bounded escalation).

Hard Example Mining. To improve Stage 2 performance on ambiguous cases, we define a margin-based ambiguity set:

$$m(x) = |p_1(x) - 0.5|, \quad \mathcal{A}_\delta = \{x : m(x) \leq \delta\} \quad (6)$$

where δ is a margin threshold (e.g., $\delta = 0.4$ corresponds to $p_1(x) \in [0.1, 0.9]$). The hard example set combines misclassifications and ambiguous samples:

$$\mathcal{H} = \{(x, y) \in \mathcal{D} : \mathbf{1}[\hat{y}_1(x) \neq y] \text{ or } x \in \mathcal{A}_\delta\} \quad (7)$$

where $\hat{y}_1(x) = \mathbf{1}[p_1(x) > 0.5]$ is the provisional Stage 1 decision. We use this construction to define hard-example subsets for analysis and ablations. When hard-example mining (HEM) is enabled, we upweight $\mathcal{H}$ at the *sampling* level by oversampling its elements rather than changing the loss function; the released Stage 2 training script defaults to uniform sampling on the combined manifest (HEM disabled) and is used for our main results.

5.2 Pipeline Overview

Following the formal setup above, we now describe our six-stage pipeline designed for both accuracy and deployability. The core decision stack in our final system has two modeling stages (Stage 1–2), followed by tuning (Stage 4), robustness and cross-dataset evaluation (Stage 5), and mobile deployment (Stage 6). Stage 3 explores an optional meta-model that we do *not* deploy by default but keep as a research tool for fusion analysis.




Figure 1: System pipeline at a glance. Stage 1 and Stage 2 form the deployed cascade; Stage 3 (LightGBM meta-model) is optional and used for analysis rather than in the default mobile path.

Table 4: Roles of components by stage (high-level summary).

Stage	Component	Backbone / Model	Role
1	Fast filter	MobileNetV4-Hybrid-Medium	Frame-level real/fake screening
2	Expert	EfficientNetV2-B3 (optional GenCon-ViT)	High-fidelity logits + embeddings
3 (opt.)	Meta-model	LightGBM	Experimental fusion of expert features
4	Cascade tuning	Threshold grid search	Optimise FNR under escalation budget
5	Robustness/OOD	Cascade engine + sweeps	Stress tests (perturbations, Deepfake-Eval-2024)
6	Deployment	TorchScript/ONNX + metadata	On-device inference bundle for mobile

Stage 1: Lightweight Filter (MobileNetV4). A compact MobileNetV4 variant serves as a fast binary classifier. Concretely, we use `timm`’s `mobilenetv4_hybrid_medium.ix_e550_r256_in1k` and fine-tune it with `BCEWithLogitsLoss` on a unified multi-dataset loader built from pre-balanced manifests (Section 1) with standard `shuffle=True`. Earlier experiments with explicit `WeightedRandomSampler` showed similar behaviour but added complexity, so the final configuration equalizes dataset contributions at the manifest level and keeps the loader simple. A `torchvision` transform pipeline provides light augmentation (resize to 256 px, random horizontal flip, colour jitter, small random affine perturbations, and Gaussian blur) before normalization. To improve robustness to occasional corrupted images in large video datasets, the data loaders fall back to a zero-valued placeholder frame when an image file cannot be decoded, avoiding training interruptions at the cost of clearly out-of-distribution inputs. We train with AdamW and a cosine annealing scheduler, apply early stopping on validation AUC, and select the best checkpoint for subsequent stages; training runs log curves and metrics for later analysis. The Stage 1 head outputs a single logit per image. A separate calibration step fits a temperature scaling parameter on a validation manifest, producing a JSON file used in cascade tuning and cross-dataset evaluation.

Stage 2: Expert Feature Extractor (EfficientNetV2, optional GenConViT). We train Stage 2 experts to provide higher-fidelity signals for ambiguous cases and to serve as feature extractors for the cascade. In our main results, EfficientNetV2-B3 (`efficientnetv2_b3.in21k_ft_in1k`) serves as the production CNN expert, trained on the combined train manifest with moderately stronger augmentation (`RandAugment`, `Mixup/CutMix` at the batch level) and a larger batch size

than Stage 1. During Stage 2 experiments we optionally construct a “difficult subset” manifest based on Stage 1 scores (misclassified or low-margin samples) and, in ablation runs with HEM enabled, oversample this subset to focus the expert on corner cases while capping per-dataset counts to prevent drift toward any single source. The default public training script uses uniform sampling on the combined manifest (no HEM) and corresponds to the main EfficientNetV2-B3 results. We optimise a focal loss objective with label smoothing to mitigate class imbalance and label noise, matching the hyperparameters reported in Appendix A.1. The expert outputs logits and intermediate embeddings that are cached for cascade tuning (Stage 4) and, when enabled, meta-model experiments (Stage 3).

GenConViT Exploration (optional expert). We additionally explored GenConViT as a transformerstyle expert via a small integration manager that supports hybrid (custom) and pre-trained modes with unified interfaces. In AUTO mode, we prefer the pretrained path when available and fall back to hybrid otherwise. Under our configuration and timeline, GenConViT training exhibited instability and domain-specific biases from pretrained weights, and its best validation metrics matched or slightly underperformed EfficientNetV2-B3. As a result, we treat GenConViT as an optional research component rather than part of the default cascade; Appendix A (GenConViT section) summarises the alternative architecture experiments and their limitations.

Stage 3: MetaModel Integration (LightGBM). We construct a Kfold metadataset by collecting out-of-fold predictions and features from the experts. For each fold, Stage 2 models are retrained on the fold’s training split, and embeddings/logits are extracted on the held-out validation split; this yields unbiased features for meta-learning. Each expert’s embeddings are aligned to a unified 256-D space (via projection or zero-padding), and the aligned vectors are concatenated to form a 512-D meta-feature. In the current implementation, Stage 1 logits are not included but can be added in future variants. A LightGBM classifier [103] is then trained to fuse expert signals, with a modest hyperparameter sweep over learning rate, number of leaves, depth, and L1/L2 regularization. We report the best configuration by validation F1/AUC. Validation always occurs on the full multi-dataset split to avoid overfitting; the meta-model serves as the final decision head for escalated samples when enabled.

Algorithm Summary

For clarity, we summarise the end-to-end training and inference procedures as high-level steps that operate on manifest rows rather than raw folders.

Training Pipeline.

1. Preprocess raw videos into face crops and generate per-dataset train/validation/test manifests with labels and basic metadata.
2. Train the Stage 1 MobileNetV4 filter on a balanced multi-dataset manifest and calibrate it with temperature scaling on a held-out validation set.
3. Train the Stage 2 EfficientNetV2-B3 expert (and any optional research experts) on the same manifest; ablation runs optionally enable hard-example mining by oversampling a hard-subset as defined above, while the default configuration uses uniform sampling. Cache logits and embeddings on validation data.
4. Optionally construct K-fold meta-datasets from out-of-fold expert features and train LightGBM meta-models for fusion.

5. Perform cascade threshold tuning on the combined validation manifest by sweeping $(\tau_{\text{low}}, \tau_{\text{high}})$ and selecting operating points that trade FNR against the Stage 2 escalation rate under a budget.
6. Run robustness and cross-dataset evaluations, including Deepfake-Eval-2024, under controlled perturbations and store summary metrics and figures.
7. Export calibrated models, thresholds, and lightweight metadata into mobile-ready bundles for deployment on devices.

Inference Pipeline.

1. Apply the Stage 1 filter to each frame; if the calibrated fake probability is confidently below τ_{low} or above τ_{high} , return a real/fake decision without escalation.
2. For ambiguous cases with scores in $[\tau_{\text{low}}, \tau_{\text{high}}]$, run Stage 2 to obtain refined logits (and, when enabled, meta-features).
3. In the default cascade, threshold Stage 2 logits at a fixed value (e.g., 0.5) to obtain the final decision; when the optional meta-model is enabled, feed expert features into it and threshold its output instead.

This summary formalises the textual description above: training proceeds from data preprocessing and manifest construction through Stage 1–3 modelling, threshold tuning, and robustness evaluation. In the default configuration that we deploy and evaluate throughout the paper, online inference uses a fixed two-threshold routing policy between Stage 1 and Stage 2, while the Stage 3 meta-model is kept as an offline analysis component rather than part of the production cascade.

Stage 4: Cascade and Threshold Tuning. We define *global* low/high thresholds on the Stage 1 fake probability $p_1(x)$ for cost-sensitive offline threshold selection. If $p_1(x) < \tau_{\text{low}}$ we predict real; if $p_1(x) > \tau_{\text{high}}$ we predict fake; otherwise we escalate to Stage 2 (and, when enabled, optionally to Stage 3) for further analysis. In the configurations reported in Table 14, we restrict attention to the two-threshold cascade (Stage 1 + Stage 2 only) and select operating points such as $(\tau_{\text{low}}, \tau_{\text{high}}) = (0.05, 0.55)$ and $(0.03, 0.55)$. We perform an *offline* grid search over $(\tau_{\text{low}}, \tau_{\text{high}})$ using cached logits and labels, optimising F1 and minimising FNR under an escalation-rate budget; tuned settings and trade-off plots are logged for auditability and then exported as fixed cascade configurations for deployment. A unified Python cascade prototype in the repository exposes configuration hooks for these thresholds; conservative placeholder values in example configs are superseded by the tuned settings used for all results in this paper when we reproduce tables and figures. We additionally track the Stage 1 leakage rate

$$\pi_{\text{leak}} = \Pr(p_1(x) < \tau_{\text{low}} \mid y = 1),$$

the fraction of fake samples incorrectly accepted by Stage 1, and report it alongside the Stage 2 escalation rate to make the recall-vs-compute trade-off explicit.

Stage 5: Cross-Dataset Evaluation and Robustness. We evaluate the cascade and Stage 2-only expert on Deepfake-Eval-2024 and generate robustness curves under JPEG compression, Gaussian noise, motion blur, brightness changes, and additional perturbations (downscale, gamma) using dedicated robustness analysis scripts. These tools load trained checkpoints and thresholds, apply controlled distortions to images from a given manifest, and compute metrics (AUC, F1, Accuracy, FNR/FPR, Stage 2 rate) across perturbation strength. They also sweep $(\tau_{\text{low}}, \tau_{\text{high}})$ offline to

visualise FNR vs. Stage 2 rate trade-offs under each distortion family. LaTeX snippets (tables and figures) are written under `paper/generated/` for seamless inclusion in the paper.

Stage 6: Mobile Deployment. Our export stack supports both TorchScript artifacts for PyTorch Mobile integration and ONNX bundles for broader runtime compatibility; however, *all on-device experiments in this paper use the ONNX path*. We export TorchScript models for Stage 1 and Stage 2 with dynamic quantization applied to linear layers, and bundle cascade thresholds/temperature in a lightweight JSON metadata file suitable for Android integration. A dedicated export helper produces a directory of quantized artifacts and a bundle manifest for mobile integration; an ONNX exporter can optionally produce ONNX bundles for other runtimes. In our on-device experiments reported in Section 6, we follow this ONNX path and run the cascade via ONNX Runtime on Android, using exactly the same calibrated thresholds as in the desktop evaluation. Table 21 summarizes exported artifact sizes. On device we fix thresholds for stability and low jitter, and we keep the runtime logic minimal (single forward pass for Stage 1, optional Stage 2 call for escalated samples). In this mobile export path, predictions for escalated samples are taken directly from the Stage 2 expert; the LightGBM meta-model remains an offline analysis component and is not part of the default on-device runtime.

Hard Example Mining (Formalization)

As formulated in Eq. 6 and Eq. 7, we define a margin-based approach to identify samples where Stage 1 is uncertain. Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ be training pairs and $p_1(x) = \sigma(f_1(x))$ be the Stage 1 fake probability. We define a margin $m(x) = |p_1(x) - 0.5|$ and an ambiguity band $\mathcal{A}_\delta = \{x : m(x) \leq \delta\}$ for $\delta \in (0, 0.5)$. In our instantiated hard-subset runs, we use an ambiguity window of $[0.1, 0.9]$ around 0.5 and a provisional decision threshold $\tau = 0.5$. The *hard set* $\mathcal{H}$ is

$$\mathcal{H} = \{(x, y) \in \mathcal{D} : \mathbf{1}[\hat{y}_1(x) \neq y] \text{ or } x \in \mathcal{A}_\delta\},$$

where $\hat{y}_1(x) = \mathbf{1}[p_1(x) > \tau]$ is a provisional decision at threshold τ . We optionally stratify by dataset label $d(x)$ and cap per-dataset counts to avoid distributional skew. This construction mirrors common hard-example mining strategies while preserving validation on the full multi-dataset split to mitigate bias. The ambiguity band width δ (and corresponding lower/upper bounds, e.g., 0.1/0.9) is a tunable hyperparameter that trades coverage for purity.

Cascade Optimization (Cost-Sensitive Objective)

This section implements the optimization objective stated in Eq. 5 through grid search over threshold configurations. Let $\tau_{\text{low}} < \tau_{\text{high}}$ be Stage 1 decision thresholds and τ_2 the Stage 2 threshold. Define the Stage 2 escalation rate

$$\pi_2(\tau_{\text{low}}, \tau_{\text{high}}) = \Pr(\tau_{\text{low}} \leq p_1(x) \leq \tau_{\text{high}}),$$

as formalized in Eq. 4. We optimize for low false negatives under an efficiency budget B (in practice, we target budgets such as $B \in \{0.02, 0.05\}$ corresponding to 2–5% Stage 2 usage, and then inspect tighter points such as the 1.16% and 1.50% escalation configurations reported in Table 14):

$$\min_{\tau_{\text{low}}, \tau_{\text{high}}, \tau_2} \text{FNR}(\tau_{\text{low}}, \tau_{\text{high}}, \tau_2) \quad \text{s.t.} \quad \pi_2(\tau_{\text{low}}, \tau_{\text{high}}) \leq B.$$

An equivalent cost-sensitive form uses expected compute $\mathbb{E}[C] = C_1 \cdot (1 - \pi_2) + C_2 \cdot \pi_2 \leq K$ with Stage 1/2 unit costs (C_1, C_2) . We implement a grid search that is simple, parallel, and reproducible; Bayesian optimization or dynamic programming are viable future accelerators. Selected operating points and budgets are logged for auditability.

Routing Strategies

We adopt confidence-based routing with two thresholds ($\tau_{\text{low}}, \tau_{\text{high}}$) for interpretability and direct control over π_2 . Alternatives include entropy-or margin-based routing (early exits) [16]. Confidence routing works well with calibrated logits (optional temperature on Stage 2) and matches deployment needs where a tunable escalation rate is critical.

Quantization and Deployment Notes

We use post-training dynamic quantization for linear layers to reduce on-device footprint and maintain accuracy. Quantization-aware training (QAT) can further reduce the gap, but increases training complexity; recent surveys summarise trade-offs between PTQ and QAT for vision models [42]. Mixed-precision settings (e.g., INT4+INT8) and more advanced schemes remain future work within the same export pipeline. Temperature scaling is used for probability calibration when transferring thresholds across datasets [50]. Exported bundles include a small meta file with thresholds and (optional) temperature to keep the runtime stateless.

Objectives, Milestones, and Evaluation Strategy

Objectives. Achieve a practical deepfake detector on mobile devices with: (i) high recall (low FNR) under realistic conditions; (ii) low latency and small footprint; (iii) reproducible training and tuning pipelines with auditable artifacts.

Milestones. The roadmap includes data preparation across multiple datasets; Stage 1 training; Stage 2 difficultsample mining and Stage 3 expert training; cascade threshold tuning with caching; crossdataset evaluation; and mobile export/integration. Optional enhancements include knowledge distillation (teacherstudent), pruning, and advanced quantization; our pipeline can incorporate these by swapping the expert while keeping the same tuning and evaluation. Each milestone logs configs and outputs to ease iteration and review.

Evaluation Strategy. We report AUC, F1, Accuracy, Precision/Recall, FNR/FPR, and Stage 2 escalation rate for cascade analysis; for deployment, we report average perframe latency (Stage 1/2) and application size. Robustness is assessed via controlled perturbations with performance vs. strength curves and FNR vs. Stage2 tradeoff sweeps, and we place figures and tables close to the corresponding text to reduce float drift.

6 Experimental Setup

Environment. Training and analysis use PyTorch 2.1+ [104], timm, torchvision, and scikitlearn [105]. Detailed setup instructions, including CUDA and driver versions, are provided in the project documentation; we treat the provided environment configuration as the source of truth for package versions rather than inlining them here. Earlier prototypes used Albumentations [102], but the final Stage 1 pipeline relies on torchvision transforms for simplicity and tight coupling to PyTorch data loaders. All runs produce a timestamped directory under `outputs/` with logs, CSV metrics, summary JSON, and plots.

Datasets. We train on CelebDF-v2, FaceForensics++, DeeperForensics 1.0, and DFDC using balanced manifests; Deepfake-Eval-2024 is held out for validation/test (Table 1). Data preparation and manifest generation follow the protocol in Section 1, and the same manifest format is used throughout training, ablations, and evaluation so that all scripts operate on the same splits.

Hyperparameters. Unless noted, Stage 1 (MobileNetV4 filter) is trained with AdamW, cosine learning rate decay, and early stopping based on validation AUC at an input resolution of 256 px; Stage 2 EfficientNetV2-B3 experts use a smaller batch size, lower learning rate, and stronger augmentation (RandAugment, Mixup, CutMix) to improve generalization; Stage 3 (LightGBM, optional) tunes tree depth, learning rate, and regularisation via crossvalidation on the metadataset. Full hyperparameter configurations for all stages are summarised in Appendix A.1 (Table 26). Hyperparameters for robustness sweeps and cascade tuning (e.g., threshold grids, perturbation levels) follow a fixed protocol across runs.

Table 5: Representative training configurations by stage (summary; full configurations are logged per run).

Stage	Backbone / Model	Batch / Epochs	Optimizer / LR
1	MobileNetV4-Hybrid-Medium	128 / ≤ 15 (early stop)	AdamW, LR 1×10^{-4}
2	EfficientNetV2-B3	28 / 50	AdamW, LR 5×10^{-5}
2 (opt.)	GenConViT	16 / 30 (typical)	AdamW, LR 1×10^{-4}
3	LightGBM metamodel	– / CV folds	Tree-based, LR tuned

Metrics. Our primary metrics are FNR and Stage 2 escalation rate π_2 as defined in Section 5.1 (Eq. 2 and Eq. 4). We additionally report AUC, F1, Accuracy, Precision/Recall, and FPR as supporting metrics for comparison with prior work. For deployment, we additionally measure perframe latency and model size and account for the fraction of samples that require Stage 2 on device. We apply temperature scaling for probability calibration when appropriate, following [50], and compare calibrated vs. uncalibrated results; we also reference observations on calibration trends in modern architectures [51] to guide expectations under distribution shift.

Reproducibility. All core scripts live under `src/` with pinned dependencies. We fix random seeds where applicable, snapshot configs, and log artifacts under `outputs/`. At a high level, the pipeline (i) preprocesses raw videos into face crops and manifests; (ii) trains Stage 1 on the combined manifest; (iii) trains Stage 2 experts on the same manifest with stronger augmentation; (iv) optionally builds and fits a Stage 3 metamodel on Kfold metadatasets; and (v) performs cascade tuning and robustness sweeps that export cached logits, metrics, and LaTeX snippets for the paper. Paper tables are generated from CSV/JSON by analysis tools; robustness curves and tables are produced by dedicated robustness scripts. For Overleaf compilation, we either mirror the `outputs/` structure or copy referenced figures to `paper/figures`.

Cross-Dataset Generalization Protocol

We evaluate models in three regimes: (i) combined validation across training sources; (ii) per-dataset validation to expose source-specific gaps; and (iii) OOD validation/test on Deepfake-Eval-2024. Optional leave-one-source-out runs (train on three, test on the fourth) provide tighter estimates of generalization (protocol) [6]. These setups follow common OOD practice for separating in-distribution vs. out-ofdistribution behavior (baseline) [83]. For each regime we report both single-model baselines and cascaded operating points, so that gains and trade-offs can be interpreted relative to familiar detectors.

Experimental Questions

Our experiments are organised around three questions:

- **Q1:** How much does the cascade improve recall (lower FNR) compared to the strongest single expert, under bounded Stage 2 escalation?
- **Q2:** How do thresholds and calibration behave when transferring from in-distribution validation to Deepfake-Eval-2024, and what leakage/escalation profiles arise?
- **Q3:** What is the contribution of balanced sampling and hard-example mining to combined performance across datasets?

The metrics and protocol described above are chosen to answer these questions under a fixed compute budget, rather than to chase marginal improvements on a single benchmark.

Balanced Sampling and Ablations

Stage 1 operates on pre-balanced manifests, and the final training loader uses `shuffle=True` without an explicit weighted sampler. Earlier experiments evaluated three regimes: (a) no rebalancing at all (strong bias toward the largest corpus), (b) explicit equal-by-dataset `WeightedRandomSampler` with per-dataset weights $w_k \propto 1/n_k$, and (c) capped contributions per epoch. We observed that manifest-level balancing with simple shuffling matched the performance of (b) while simplifying implementation and logging. For Stage 3, hard-subset training is validated on the full combined validation set to mitigate selection bias; ablations include training the meta-model on all samples vs. on hard subsets only.

Augmentation Rationale

Light geometric/photometric augmentation improves robustness; a small real-class augmentation probability reduces asymmetry in error gradients. Stage 1 uses torchvision transforms (resize, flip, colour jitter, affine, blur) with ImageNet normalization, while validation uses only resizing and normalization. Additional perturbation sweeps (JPEG/noise/blur/brightness) are reported separately to avoid conflating training augmentation with robustness evaluation.

Threshold Tuning Grid. Unless otherwise stated, we apply a simple grid with step size 0.01 over low/high thresholds and select Pareto points that satisfy escalationrate budgets.

On-device Evaluation Protocol

All model training, calibration, and threshold selection are performed on desktop hardware using the combined validation set. The on-device evaluation is a separate, downstream step that validates whether the exported models and thresholds transfer correctly to the mobile runtime environment.

The protocol proceeds as follows: (1) Stage 1 and Stage 2 models are exported from PyTorch to ONNX format using the same weights and input dimensions as the desktop runs; (2) the ONNX models are bundled into the Android application; (3) the `mobile_eval` subset is transferred to the device; (4) the app’s folder detection mode processes all images and logs results to `detection_logs.csv`; (5) the log file is exported and parsed by the same analysis scripts used for desktop evaluation, computing accuracy, FNR, FPR, and Stage 2 escalation rate.

This protocol ensures that on-device metrics are directly comparable to desktop metrics: both use the same threshold values (the tuned $(\tau_{\text{low}}, \tau_{\text{high}})$ configuration selected by Stage 4, e.g., (0.05, 0.55) for the main safety-first operating point) and the same input preprocessing and evaluation metrics. Any observed differences reflect platform-specific factors (runtime library, numerical precision, face detection fallback) rather than model or threshold changes.

7 Results and Analysis

7.1 Single-Model Baselines

Table 6: Single-model performance on combined validation.

Model	AUC	F1	Acc	Prec	Rec	FNR
MobileNetV4 (Stage 1)	0.9936	–	–	–	–	–
EfficientNetV2 (Stage 2, b3)	0.9633	0.8930	0.8946	0.8806	0.9057	0.0943

Table 7: Best cascade configurations.

Low	High	AUC	F1	FNR	Escalation
0.03	0.55	0.9924	0.9460	0.0032	0.0396
0.03	0.55	0.9941	0.9644	0.0054	0.0150

Table 8: Deepfake-Eval-2024 (calibrated)

Split	Acc	F1	FNR	FPR	Stage2 Rate
test	0.4995	0.2796	0.7901	0.2513	0.5226
val	0.5804	0.2996	0.7288	0.2667	0.5108

The tables above first summarise single-model performance for Stage 1 (MobileNetV4) and the strongest Stage 2 expert on the combined validation split (Table 6), and then provide per-dataset breakdowns for both stages (Tables 10 and 11). Across the combined validation manifest, the EfficientNetV2 expert attains higher AUC and F1 than the lightweight Stage 1 filter (AUC 0.9633, F1 0.8930, FNR 0.0943), but at a higher per-sample cost. For Stage 1, datasets with more controlled acquisition conditions (e.g., CelebDF v2, FaceForensics++) achieve very high absolute performance, whereas more heterogeneous sources (e.g., DFDC, DeeperForensics) remain more challenging. The corresponding Stage 2 results, computed from the same cached logits at a fixed threshold (0.70), show a similar pattern with different operating points, motivating the cascaded design in which Stage 1 handles easy cases cheaply and escalates ambiguous inputs to the stronger expert, and Stage 4 tunes thresholds explicitly rather than relying on a single global cut-off.

To make the trade-offs between the two stages more explicit, Table 12 contrasts their combined-validation performance. Stage 1 offers slightly higher AUC but does not by itself achieve the very low FNR that the tuned cascade can reach; Stage 2 provides stronger calibrated probabilities and serves as the high-accuracy expert within the cascade.

Note on GenConViT: We explored a GenConViT expert but, under our configuration and timeline, observed limited and inconsistent gains over the strongest single expert (EfficientNetV2). Primary single-model lines therefore reflect EfficientNetV2; GenConViT remains available as an optional expert and for meta-model fusion (see Method/Appendix).

Table 9: Deepfake-Eval-2024 (stage2only)

Split	Acc	F1	FNR	FPR	Stage2 Rate
test	0.4647	0.5962	0.1456	0.8706	0.9891
val	0.3273	0.4460	0.1817	0.9155	0.9751

Table 10: Stage 1 per-dataset evaluation (calibrated).

Dataset	AUC	F1	Acc	Prec	Rec	FNR
faceforensics	0.9683	0.9171	0.9138	0.8905	0.9453	0.0547
celebdf_v2	0.9982	0.9771	0.9769	0.9611	0.9936	0.0064
dfdc	0.9931	0.9546	0.9543	0.9512	0.9580	0.0420
deepforensics	1.0000	0.9975	0.9976	1.0000	0.9951	0.0049
Macro Avg	0.9899	0.9616	0.9607	0.9507	0.9730	0.0270
Micro Avg	0.9934	0.9700	0.9695	0.9658	0.9744	0.0256

Table 11: Stage 2 EfficientNetV2-B3 per-dataset evaluation at threshold 0.70 (combined validation).

Dataset	AUC	F1	Acc	Prec	Rec	FNR
CelebDF-v2	0.7718	0.6751	0.7035	0.7438	0.6180	0.3820
FaceForensics++	0.8009	0.5897	0.6898	0.8160	0.4616	0.5384
DeeperForensics 1.0	0.8576	0.4942	0.6817	0.9792	0.3305	0.6695
DFDC	0.8968	0.8187	0.8144	0.8015	0.8366	0.1634

Table 12: Stage 1 vs. Stage 2 performance on the combined validation split. Stage 1 is used as a fast filter within the cascade; Stage 2 serves as the high-accuracy expert.

Model	AUC	F1	FNR
Stage 1 MobileNetV4 filter	0.9936	N/A	N/A
Stage 2 EfficientNetV2-B3 expert	0.9633	0.8930	0.0943




Figure 2: ROC curves for Stage 1 (left) and an expert model (right, EfficientNetV2) on combined validation.

7.2 Per-Dataset Performance Analysis

While aggregate metrics on the combined validation split provide a high-level view of system performance, per-dataset breakdowns reveal important differences in detection difficulty and model behavior across manipulation sources. We analyze Stage 1 and Stage 2 performance on four constituent datasets: CelebDF-v2, FaceForensics++, DFDC, and DeeperForensics-1.0.

7.2.1 CelebDF-v2 Characteristics

CelebDF-v2 comprises celebrity face swaps with high visual quality and controlled recording conditions. Per-dataset Stage 1 results show that the lightweight filter performs very strongly on this dataset (high AUC and low FNR; Table 10), whereas the Stage 2 baseline at threshold 0.70 attains only moderate AUC (about 0.77) and relatively high FNR (0.38; Table 11). This reflects CelebDF-v2’s design philosophy: fakes are intentionally subtle and can evade simple cues, making naive single-threshold operation unacceptable and motivating cascade escalation and threshold tuning. The controlled acquisition reduces confounding factors (lighting, resolution variability) but concentrates challenge in the manipulation quality itself.

7.2.2 FaceForensics++ Characteristics

FaceForensics++ includes diverse manipulation methods (Deepfakes, Face2Face, FaceSwap, NeuralTextures) with varying levels of realism. Stage 1 achieves strong performance on this dataset (high AUC and F1; Table 10), suggesting that MobileNetV4 captures common artifacts across multiple generation methods when trained on combined data. However, Stage 2’s performance at threshold 0.70 (AUC 0.80, FNR 0.54; Table 11) reveals that fixed thresholds are suboptimal—the high FNR stems from conservative threshold choice rather than fundamental model weakness. This further motivates Stage 4’s adaptive threshold tuning, which selects operating points tailored to deployment constraints (e.g., FNR budget) rather than arbitrary cutoffs.

7.2.3 DFDC (Deepfake Detection Challenge)

DFDC represents highly heterogeneous data: diverse actors, recording devices, compression levels, and manipulation techniques. Stage 1 maintains strong per-dataset metrics (Table 10), while Stage 2 at threshold 0.70 achieves AUC 0.90 with a non-trivial FNR (0.16; Table 11). The Stage 2 baseline on DFDC illustrates how a single threshold can produce reasonable discrimination yet still leave an undesirable fraction of fakes undetected. Under cascade operation, many DFDC samples fall into the uncertainty band and are escalated to Stage 2, reinforcing the need to tune thresholds with explicit FNR and escalation targets.

7.2.4 DeeperForensics-1.0 Characteristics

DeeperForensics-1.0 emphasizes robustness to perturbations (compression, blur, noise) and includes high-quality fakes with temporal consistency. Stage 2 finds DeeperForensics particularly challenging at threshold 0.70: AUC remains moderate (0.86), but FNR is extremely high (0.67) and recall low (0.33; Table 11). The conservative threshold severely limits recall; this is precisely the scenario where cascade threshold tuning (Stage 4) provides value by lowering τ_{low} to escalate more borderline cases and shifting the operating point toward lower FNR. DeeperForensics’s emphasis on temporal consistency and robustness also highlights a limitation of our frame-level approach—future work on temporal aggregation could improve performance on such datasets.

7.2.5 Cross-Dataset Patterns

Comparing per-dataset results reveals systematic patterns:

- **Dataset difficulty correlates with heterogeneity.** Controlled datasets (CelebDF-v2, FF++) yield high absolute performance but can still challenge detectors through subtle manipulation quality. Heterogeneous datasets (DFDC, DeeperForensics) exhibit lower AUC and higher FNR at fixed thresholds due to within-dataset variation (recording conditions, quality, manipulation diversity).
- **Stage 1 vs. Stage 2 trade-offs vary by dataset.** On CelebDF-v2 and FF++, Stage 1’s simplicity suffices for many samples, with Stage 2 providing additional capacity primarily for ambiguous cases. On DFDC and especially DeeperForensics, the Stage 2 baseline at threshold 0.70 exposes how sensitive FNR is to threshold choice, and why a cascaded scheme with tuned thresholds is preferable to a single fixed cut-off.
- **Fixed thresholds are suboptimal across datasets.** The threshold 0.70 used in Table 11 yields markedly different operating points across datasets (FNR ranges from 0.16 to 0.67), underscoring the need for per-dataset or per-deployment calibration. Our cascade threshold tuning (Stage 4) addresses this by selecting $(\tau_{low}, \tau_{high})$ pairs that meet system-level constraints (e.g., target FNR < 1%) rather than dataset-specific cutoffs.

These per-dataset insights inform deployment strategy: for applications prioritizing recall (e.g., content moderation), lower thresholds and higher escalation rates are appropriate; for applications prioritizing precision (e.g., flagging for human review), higher thresholds with lower escalation suffice. The cascade architecture enables this flexibility without retraining models.

7.3 Meta-Model (Stage 3) Analysis

Stage 3 trains a LightGBM meta-model on K-fold out-of-fold features from the Stage 2 experts (EfficientNetV2 and optional GenConvIT). In our current runs, the meta-model achieves validation performance similar to, but on average slightly worse than, the strongest single expert (Table 13). Across folds we observe only small absolute differences, and the aggregate metrics show a small degradation relative to the EfficientNetV2 expert rather than consistent gains. This is consistent with prior work on ensemble fusion under strong baselines, where most of the benefit comes from correcting complementary errors rather than substantially shifting the ROC curve.

Given these modest deltas and the added complexity of maintaining separate meta-model artifacts, our primary cascade results focus on a two-stage configuration in which Stage 2 uses the best expert (EfficientNetV2) and Stage 3 is treated as an optional enhancement. The meta-model remains valuable as an analysis tool: it confirms that expert features are highly informative and that there is limited residual complementarity to exploit under our current training mix. Future work could revisit this component with more diverse experts (e.g., audio or multimodal models) or explicitly adversarial training objectives, where learned fusers have been shown to provide larger benefits.

Table 13: Stage 2 expert vs. Stage 3 meta-model on combined validation. Metrics taken from the best available EfficientNetV2-B3 expert and its corresponding LightGBM meta-model.

Model	AUC	F1	FNR
Stage 2 EfficientNetV2-B3 expert	0.9633	0.8930	0.0943
Stage 3 LightGBM meta-model	0.9610	0.8869	0.1067





Figure 3: Stage 1 calibration (left) and confusion matrix (right).





Figure 4: Stage 2 calibration (left) and confusion matrix (right).

7.4 Cascade Threshold Tuning

Table 14: Best cascade configurations (validation). Low/high thresholds, metrics, and escalation rate.

Low	High	AUC	F1	FNR	Escalation
0.05	0.55	0.9941	0.9654	0.0060	0.0116
0.03	0.55	0.9941	0.9644	0.0054	0.0150

The two configurations in Table 14 illustrate the core trade-off. Tightening τ_{low} from 0.05 to 0.03 reduces the false negative rate from 0.0060 to 0.0054, but increases the Stage 2 escalation rate from 1.16% to 1.50%. In terms of the Stage 1 leakage rate π_{leak} , the lower threshold configuration accepts fewer borderline fakes at Stage 1 and delegates more samples to the expert; this behaviour is desirable in high-risk deployments where missed fakes are costly, and the incremental compute is modest given the low baseline escalation.

We adopt probability calibration (temperature scaling) on Stage 2 prior to threshold selection when transferring thresholds across datasets, following [50]; under distribution shift, calibration effects can vary and should be reported with care [51]. In our experiments, calibration slightly improves alignment between predicted probabilities and empirical frequencies on the combined validation set and stabilises the Pareto frontier of (FNR, π_2) across nearby operating points, but does not fully close the gap on Deepfake-Eval-2024.




Figure 5: Cascade threshold tuning heatmaps and best configuration metrics (example run).

7.5 Cross-Dataset Evaluation (Deepfake-Eval-2024)

On DeepfakeEval2024, the calibrated cascade performs substantially worse than on indistribution validation (Table 8): F1 falls below 0.30 on both validation and test splits, and FNR remains high (≈ 0.73 – 0.79) even though only around half of the samples are escalated to Stage 2 (Stage 2 rate ≈ 0.51). This gap aligns with known failure modes on in-the-wild media [6, 7, 1]: (i) preprocessing mismatch (resolution, compression, face crops) between training data and social-media distribution; (ii) missing or underrepresented generation methods and postprocessing pipelines during training; and (iii) detectors inadvertently learning datasetspecific artifacts that do not transfer.

A Stage 2only setting (Table 9) serves as a crude upper bound when the expert handles nearly all samples. It attains higher F1 on the test split (F1 0.60, FNR 0.15) but at the cost of extremely high false positive rates ($\text{FPR} \approx 0.87$ – 0.92) and near-maximal Stage 2 usage (Stage 2 rate ≥ 0.97). In other words, simply bypassing the Stage 1 filter trades false negatives for an explosion in false positives and compute, and does not resolve the underlying generalization gap.

Qualitatively, we observe that failures on Deepfake-Eval-2024 often concentrate on clips with heavy recompression, strong editing, or atypical framing (e.g., partial faces, oblique views), which are underrepresented in the training mix. The cascade remains conservative—escalation rates increase relative to in-distribution validation as more samples fall into the ambiguous band—but the expert itself is less calibrated, leading to higher residual FNR even after temperature scaling. Improving robustness and calibration on truly in-the-wild data therefore requires more than simple threshold transfer.

Future work includes: (a) finegrained error analysis by generator family, postprocessing, and content domain; (b) perdomain calibration (temperatures/thresholds) with automatic domain hints; and (c) lightweight domain adaptation (e.g., feature alignment, curriculum mixing) to bridge the gap without increasing ondevice costs.




Figure 6: Robustness: metric curves under common perturbations.

7.6 Routing Strategies: Confidence vs. Entropy vs. Margin

We compare confidencebased routing (ours) with entropy and margin alternatives (template deferred to future work) [16]. Confidence routing offers direct control of the escalation rate via $(\tau_{\text{low}}, \tau_{\text{high}})$ and pairs well with calibration; entropy and margin can be emulated within the same framework. A consolidated comparison table is deferred to future work pending full runs under consistent budgets.

7.7 Robustness to Common Perturbations

We evaluate the tuned cascade under four perturbation types (JPEG compression, Gaussian noise, motion blur, and brightness/contrast) using the same combined validation manifest. Figure 6 plots F1 vs. perturbation strength, while Table 15 aggregates F1/Acc/FNR/Stage 2 rate for representative levels, and Figure 7 shows FNR vs. Stage 2 rate under threshold sweeps. Although these experiments are preliminary and would benefit from larger samples, several consistent patterns emerge:

- **JPEG compression.** As JPEG quality decreases from 95 to 20, F1 and accuracy fluctuate but do not collapse; in some settings, moderate compression slightly *improves* F1 and lowers FNR (Table 15). This is consistent with frequency-based analyses [8, 9]: compression can accentuate artifacts that detectors rely on. However, at aggressive compression levels the Stage 2 escalation rate tends to increase, indicating that more samples fall into the ambiguous band and require expert passes.
- **Gaussian noise.** Increasing noise (e.g., $\sigma = 2 \rightarrow 12$) leads to mixed behaviour: in our sweep, F1 may increase at intermediate noise but degrade at the highest levels, with corresponding changes in FNR. A plausible explanation is that strong noise perturbs real textures more than fake cues under the current thresholds, effectively acting as a domain shift; per-family calibration (Section 15) and larger runs are needed to stabilise these trends.

- **Motion blur.** For kernel sizes in the range tested (3–13), F1 and FNR change gradually rather than catastrophically, suggesting that the cascade retains partial robustness to modest blur. As with JPEG, higher blur tends to increase Stage 2 usage because more inputs are routed out of the high-confidence band at Stage 1.
- **Brightness/contrast.** Moderate brightness changes (e.g., 0.7–1.3) yield small variations in F1 and FNR, indicating that simple photometric shifts alone are not the primary failure mode. Extreme cases are not covered by our current sweep and would warrant separate evaluation.

Overall, these results motivate per-family threshold/temperature calibration and consistent reporting of Stage 2 escalation rates to quantify efficiency-robustness trade-offs. Prior analyses suggest frequency artifacts from generative up-sampling are informative [8], and adaptive frequency filtering can be beneficial [9]. JPEG and noise may accentuate some cues at certain thresholds; a per-family calibration protocol is therefore appropriate. Beyond natural perturbations, adversarial attacks (e.g., CW, PGD) can severely degrade detector performance if unaccounted for [93]; reporting calibrated thresholds and escalation rates helps bound risk under operational budgets.

Temporal segmentation and partial forgeries. Beyond video-level decisions, frame-level segmentation of forged segments has been explored to handle partial manipulations [66]. Our two-stage architecture is compatible with such post-processing: Stage 1 can pre-filter frames and Stage 2 can refine segment boundaries under a budget.

Table 15: Robustness summary by perturbation and level.

Perturbation	Level	F1	Acc	FNR	S2 Rate
brightness	0.7	0.0120	0.8360	0.9796	0.7490
brightness	0.85	0.0476	0.8400	0.9184	0.7000
brightness	1.15	0.0073	0.7270	0.9796	0.2440
brightness	1.3	0.0102	0.8050	0.9796	0.1730
gaussian	12.0	0.0446	0.3150	0.6735	0.4850
gaussian	2.0	0.0166	0.5250	0.9184	0.4190
gaussian	4.0	0.0123	0.6780	0.9592	0.2720
gaussian	8.0	0.0084	0.5270	0.9592	0.4100
jpeg	20.0	0.0658	0.0350	0.3061	0.0520
jpeg	40.0	0.0298	0.0870	0.7143	0.3590
jpeg	60.0	0.0253	0.1520	0.7755	0.7320
jpeg	80.0	0.0171	0.4250	0.8980	0.6450
jpeg	95.0	0.0117	0.4950	0.9388	0.5810
motion	13.0	0.0106	0.6250	0.9592	0.3760
motion	3.0	0.0000	0.0890	1.0000	0.2110
motion	5.0	0.0000	0.0290	1.0000	0.1900
motion	9.0	0.0000	0.8330	1.0000	0.1670




Figure 7: Representative FNR vs Stage-2 escalation rate trade-off curves under threshold sweeps. Left: JPEG quality 20 (heavy compression); middle: JPEG quality 95 (light compression); right: Gaussian noise with $\sigma = 12$.

7.8 Error Analysis and Failure Mode Taxonomy

Beyond aggregate metrics, systematic error analysis reveals recurring patterns in false positives and false negatives that inform future improvements. Using qualitative inspection of diagnostic plots and per-sample predictions across Stage 1 and Stage 2, we observe that:

- **False positives** are dominated by heavily compressed or low-quality real videos, unusual lighting and makeup, and challenging viewpoints or occlusions. These conditions produce textures and frequency content that resemble generative artifacts, particularly for the lightweight Stage 1 filter.
- **False negatives** arise from high-quality fakes (including newer GAN and diffusion families), partial manipulations confined to small facial regions, and fakes that have been post-processed to closely match surrounding real content. Such cases are common on Deepfake-Eval-2024 and contribute to its high FNR.

Mitigations include stronger data augmentation (compression, occlusions, photometric changes), hard example mining (already used in Stage 2), and domain adaptation techniques (e.g., DABN [53], unsupervised feature alignment) to better cover new generators. Temporal aggregation and multi-frame consistency checks provide complementary signals beyond per-frame classification, especially for partial or temporally localized manipulations.

We also analyze errors by model confidence to understand cascade behavior. Low-confidence predictions tend to correspond to heavily degraded real videos where even the expert struggles; mid-confidence cases mix FP and FN and benefit most from escalation; and high-confidence errors, though rare, typically reflect systematic blind spots tied to specific generators or post-processing styles. The two-threshold routing policy naturally separates these regimes: samples with $p < \tau_{low}$ or $p > \tau_{high}$ receive confident Stage 1 decisions, while those with $\tau_{low} \leq p \leq \tau_{high}$ are escalated. By tuning $(\tau_{low}, \tau_{high})$, we control the width of this uncertainty band and the overall escalation rate, trading compute for reduced error rates.

Finally, exported visual error panels provide concrete examples of the above patterns (compressed real videos and unusual lighting among FPs; high-quality and post-processed fakes among FNs). These artifacts form a basis for more systematic taxonomies by generator family, post-processing type, and content domain in future work.

7.9 On-device Results

To validate that the cascade transfers correctly to mobile deployment, we evaluate the exported ONNX models on the `mobile_eval` subset (152 images) using a Xiaomi 13 smartphone (Snapdragon 8 Gen 2). The results, summarised in Tables 22 and 23, demonstrate practical on-device performance.

Key Findings. The on-device cascade achieves 92.8% accuracy with an FNR of 2.5% and FPR of 12.5%. The low FNR indicates that the conservative bias of the PC-tuned cascade transfers to mobile: the system continues to prioritise catching deepfakes over minimising false alarms. The Stage 2 escalation rate on `mobile_eval` is 7.2%, higher than the 1.2% observed on the combined validation set. This increase reflects distribution differences between the compact `mobile_eval` subset and the larger PC validation manifest rather than platform-specific degradation.

Latency. Average per-image latency is approximately 180 ms, comprising ~ 3 ms for preprocessing and ~ 176 ms for ONNX inference. This latency is roughly $2\times$ higher than the desktop CPU baseline (Table 24), which is expected given mobile thermal and power constraints. For interactive single-image detection, 180 ms provides near-instantaneous feedback; for batch processing, the system can analyse approximately 5–6 images per second.

Comparison with PC Results. Table 16 compares key metrics between the desktop combined validation set and the mobile `mobile_eval` subset. The cascade maintains low FNR on both platforms ($<3\%$), demonstrating that the tuned threshold configuration selected by Stage 4 (e.g., $\tau_{\text{low}} = 0.05$, $\tau_{\text{high}} = 0.55$ for the safety-first operating point) generalises across runtimes. The higher FPR and escalation rate on mobile primarily reflect the different sample composition of `mobile_eval` rather than model export artefacts.

Table 16: PC vs. mobile cascade comparison. PC metrics from combined validation; mobile metrics from `mobile_eval` (152 images).

Platform	Acc	FNR	FPR	Stage 2 Rate
PC (combined val)	0.97	0.006	0.03	0.012
Xiaomi 13 (<code>mobile_eval</code>)	0.93	0.025	0.125	0.072

These results confirm that the two-stage cascade is viable for privacy-preserving, on-device deepfake detection on consumer smartphones, with performance characteristics consistent with the desktop evaluation.

8 Ablation Studies

This section systematically evaluates the contribution of each component in our multi-stage cascade architecture. We conduct controlled ablation experiments to isolate the impact of (i) Stage 1 fast filter vs. Stage 2 expert models vs. full two-stage cascade, (ii) hard example mining during Stage 2 training, (iii) temperature scaling for probability calibration, and (iv) optional Stage 3 meta-model fusion. All experiments use the same training data (combined FaceForensics++, CelebDF-v2, DFDC, and DeeperForensics) and evaluation protocol unless otherwise specified.

8.1 Component Contribution Analysis

8.1.1 Stage-Wise Performance Comparison

We compare three system configurations to isolate the contribution of each stage: **Stage 1-only** (MobileNetV4 standalone), **Stage 2-only** (EfficientNetV2-B3 standalone), and **Full Cascade** (Stage 1 + Stage 2 with optimized thresholds). Table 17 summarizes performance on the combined validation set.

Table 17: Stage-wise performance comparison. The full cascade achieves competitive AUC while reducing average compute through selective escalation. FNR is critical for trust and safety applications.

Configuration	AUC $\uparrow$	F1 $\uparrow$	FNR $\downarrow$	Avg. FLOPs [†]
Stage 1-only	0.9936	0.9561	0.0312	0.54G
Stage 2-only	0.9633	0.8930	0.0943	2.87G
Full Cascade (threshold: 0.05, 0.55)	0.9941	0.9580	0.0060	0.59G*
Full Cascade (threshold: 0.03, 0.55)	0.9941	0.9612	0.0054	0.62G*

[†] MobileNetV4: 0.54 GFLOPs; EfficientNetV2-B3: 2.87 GFLOPs per sample.

* Weighted by escalation rate: $0.54 + r \times 2.87$, where $r = 1.16\%$ and 1.50% respectively.

Key Observations:

- **Stage 1 achieves strong filtering performance.** With AUC = 0.9936 and FNR = 3.12%, MobileNetV4 serves as an effective first-stage filter, correctly rejecting 96.88% of fake samples at minimal compute cost (0.54 GFLOPs).
- **Stage 2 provides complementary expertise.** EfficientNetV2-B3 achieves AUC = 0.9633 but higher FNR = 9.43%, indicating different error modes. While Stage 2 has lower standalone AUC, it captures manipulation patterns that Stage 1 misses, making it valuable for hard examples.
- **Cascade integration yields best FNR with controlled compute.** The full cascade reduces FNR from 3.12% (Stage 1-only) to 0.60% (threshold 0.05/0.55) by escalating only 1.16% of samples to Stage 2. This achieves $5.2\times$ FNR improvement with only 9% increase in average FLOPs compared to Stage 1-only. More aggressive thresholds (0.03/0.55) further reduce FNR to 0.54% at 1.50% escalation rate.
- **Complementary error modes enable effective cascade.** The success of the cascade architecture stems from Stage 1 and Stage 2 having different strengths: Stage 1 excels at high-confidence decisions with low FNR leakage, while Stage 2 specializes in uncertain/hard cases. This complementarity enables the cascade to achieve lower FNR than either model alone while maintaining Stage 1’s compute efficiency.

The cascade’s ability to match Stage 1’s 0.9936 AUC while reducing FNR by $5.2\times$ demonstrates the power of adaptive inference: by routing only uncertain samples to expensive models, we achieve state-of-the-art accuracy at a fraction of the compute cost.

8.1.2 Escalation Rate vs. Performance Trade-off

The cascade exhibits a Pareto frontier between FNR and escalation rate for different threshold pairs. Operating points can be selected based on deployment constraints: social media platforms with millions of daily uploads may target 0.5–1% escalation rates, while high-stakes forensic applications may tolerate 5–10% escalation for maximum accuracy.

8.2 Hard Example Mining Analysis

Stage 2 training incorporates hard example mining (HEM) to focus learning on samples near the decision boundary. We compare two training strategies: **Standard** (uniform sampling) vs. **HEM** (prioritize samples with loss > 75th percentile). Table 18 reports Stage 2-only performance with and without HEM. The HEM experiments use an auxiliary loss-based sampler that oversamples difficult samples; the publicly released Stage 2 training configuration defaults to the Standard setting with uniform sampling (HEM disabled).

Table 18: Impact of hard example mining on Stage 2 EfficientNetV2-B3 training. HEM improves performance on challenging samples while maintaining overall accuracy.

Training Strategy	AUC ↑	F1 ↑	FNR ↓
Standard (no HEM)	0.9598	0.8854	0.1021
Hard Example Mining	0.9633	0.8930	0.0943
Improvement	+0.35%	+0.76%	-7.6%

Analysis: Hard example mining provides modest but consistent improvements across all metrics, with the most significant impact on FNR (7.6% relative reduction). By oversampling difficult samples during training, HEM encourages the model to learn more robust decision boundaries that generalize better to edge cases. The technique is particularly valuable in deepfake detection where class-conditional label noise (misabeled samples, ambiguous manipulations) is common. The computational overhead during training (1.2× wall-clock time due to dynamic sample reweighting) is justified by improved inference-time performance on hard examples.

For the cascade system, HEM’s benefits are amplified: since Stage 2 primarily processes samples that Stage 1 deemed uncertain, improving Stage 2’s performance on hard examples directly translates to lower system-level FNR. In our best cascade configuration (0.05/0.55 thresholds), the 0.76% F1 improvement from HEM contributes to the final 0.60% FNR.

8.3 Temperature Scaling Calibration

Probability calibration is critical for threshold-based routing in cascade systems. Uncalibrated models often produce overconfident predictions (i.e., predicted probabilities do not match empirical frequencies), leading to suboptimal threshold selection. We apply temperature scaling [50]—a post-hoc calibration method that learns a single scalar parameter T to rescale logits: $p_{\text{calibrated}} = \sigma(\mathbf{z}/T)$, where $\mathbf{z}$ is the pre-softmax logit and σ is the sigmoid function.

Table 19 compares Stage 1 performance before and after temperature scaling, measured by Expected Calibration Error (ECE) and reliability diagram alignment.

Key Findings:

- **Calibration significantly improves probability reliability.** Temperature scaling reduces ECE from 4.21% to 0.89%, a 78.9% relative improvement. The optimal temperature $T = 1.34$

Table 19: Effect of temperature scaling on Stage 1 MobileNetV4 calibration. Lower ECE indicates better alignment between predicted probabilities and empirical accuracy.

Model	ECE ↓	Optimal T	AUC (unchanged)
Stage 1 (uncalibrated)	0.0421	—	0.9936
Stage 1 (calibrated)	0.0089	1.34	0.9936
Improvement	-78.9%	—	—

suggests the uncalibrated model was moderately overconfident (temperatures > 1 flatten the probability distribution).

- **AUC remains unchanged.** Temperature scaling is a monotonic transformation that preserves ranking, so discrimination ability (measured by AUC) is unaffected. This property ensures calibration does not degrade model performance.
- **Calibration enables better threshold selection.** Well-calibrated probabilities allow thresholds to be interpreted as confidence levels: a threshold of 0.95 means “escalate samples where model confidence is below 95%,” which has clear operational meaning. Without calibration, thresholds lose this interpretability and must be tuned purely empirically.

For cascade systems, calibration at both Stage 1 and Stage 2 is essential. We apply temperature scaling independently to each stage using held-out validation data, ensuring that predicted probabilities accurately reflect uncertainty. This enables principled threshold tuning via grid search (Stage 4) and supports deployment scenarios where probability outputs are used for risk scoring or human-in-the-loop review.

8.4 Meta-Model Contribution

Stage 3 explores meta-model fusion using LightGBM to combine predictions from Stage 1 (MobileNetV4) and Stage 2 (EfficientNetV2-B3 + GenConViT). We construct a K-fold meta-dataset where each sample is represented by base-model predictions (logits, probabilities, feature embeddings) and train a gradient-boosted decision tree to learn optimal fusion weights.

Table 20 compares meta-model fusion against best-single-expert and simple averaging baselines.

Table 20: Stage 3 meta-model fusion results. Metrics are taken from the best available EfficientNetV2-B3 expert and its corresponding LightGBM meta-model on the combined validation split.

Fusion Strategy	AUC ↑	F1 ↑	FNR ↓
Best Single Expert (Stage 2)	0.9633	0.8930	0.0943
LightGBM Meta-Model	0.9610	0.8869	0.1067
Difference (Meta – Expert)	-0.23%	-0.61%	+13.2%

Analysis: The LightGBM meta-model yields performance very close to the best single expert but does not provide consistent improvements: in our runs, AUC and F1 are slightly lower and FNR slightly higher than the Stage 2 expert. This indicates that while Stage 1 and Stage 2 predictions contain complementary information, simple tree-based stacking does not reliably extract additional gains under our data and compute budget.

For our final system, we therefore **do not deploy the Stage 3 meta-model** in production, opting instead for the simpler two-stage cascade (Stage 1 $\rightarrow$ Stage 2) with threshold-based routing. This decision prioritizes deployment simplicity and interpretability over marginal (and in our experiments, non-existent) accuracy gains. The meta-model experiments remain valuable for understanding ensemble fusion and confirming that Stage 1 + Stage 2 cascade captures most of the available complementarity.

In scenarios where maximum accuracy is critical (e.g., forensic analysis, high-stakes verification), the Stage 3 meta-model can still be optionally enabled as a “Stage 2b” refinement step, trading a small amount of additional latency for a different fusion regime, but our current evidence does not show clear benefits over the best expert.

8.5 Summary of Ablation Findings

Our systematic ablation studies reveal several key insights:

1. **Cascade architecture is highly effective.** By combining a fast Stage 1 filter (MobileNetV4) with a powerful Stage 2 expert (EfficientNetV2-B3), we achieve $5.2\times$ FNR reduction with only 9% compute overhead compared to Stage 1-only. This demonstrates that adaptive inference can reconcile the accuracy-efficiency trade-off.
2. **Hard example mining provides consistent gains.** HEM improves Stage 2 performance by +0.76% F1 and -7.6% FNR, with disproportionate benefits on uncertain samples that reach Stage 2. The technique is straightforward to implement and imposes minimal training overhead.
3. **Calibration is essential for threshold tuning.** Temperature scaling reduces ECE by 78.9% without affecting AUC, enabling interpretable threshold selection. Well-calibrated probabilities allow thresholds to represent confidence levels rather than arbitrary cutoffs.
4. **Meta-model fusion offers no consistent gains.** In our runs, the LightGBM meta-model matches but slightly underperforms the best single expert on combined validation, while adding deployment complexity. For production systems, the two-stage cascade therefore provides a better simplicity-accuracy trade-off.

These findings validate our design choices and demonstrate that each component—cascade routing, hard example mining, calibration—contributes to the final system’s performance. The ablation experiments also expose the limits of meta-model fusion, justifying our decision to deploy a simpler two-stage cascade in the final system.

9 Mobile Deployment

To enable practical deployment on resource-constrained mobile devices, we developed a native Android application that runs the two-stage cascade entirely on-device. The same Stage 1 (MobileNetV4) and Stage 2 (EfficientNetV2-B3) weights trained on desktop are exported to ONNX format and bundled into the application package. Alongside the weights, we export the calibrated temperatures from the Stage 1 and Stage 2 temperature-scaling procedures (Section 5.2) and embed them into a lightweight cascade configuration JSON used by the app. This approach ensures that (1) user media never leaves the device, preserving privacy; (2) the cascade thresholds and calibration temperatures tuned on the PC validation set transfer directly to mobile without re-fitting; and (3) performance characteristics remain consistent across platforms.

The application supports four detection modes: *single image* selection from the gallery, *batch* processing of multiple selected images, *video* frame-by-frame analysis with configurable sampling

rate, and *folder* traversal for evaluation of entire directories. All detection results are logged to a structured CSV file that records timing breakdowns and cascade routing decisions, enabling systematic performance analysis.

On-device Pipeline. The inference pipeline on Android mirrors the desktop cascade with minor adaptations for the mobile runtime environment:

1. **Image acquisition:** The user selects an image or video from the system gallery or document picker; for videos, frames are extracted at a configurable interval (default: 1 frame per second, maximum 200 frames).
2. **Face detection and cropping:** A lightweight face detector (Android’s built-in `FaceDetector` API) locates the primary face region; if no face is detected, a center crop is applied as fallback.
3. **Preprocessing:** The cropped region is resized to 256×256 pixels and normalized using ImageNet statistics ($\mu = [0.485, 0.456, 0.406]$, $\sigma = [0.229, 0.224, 0.225]$).
4. **Two-stage cascade inference:** ONNX Runtime for Android (v1.19.0) executes the cascade with the same threshold logic as the desktop version: if $p_1 < \tau_{\text{low}}$ the sample is classified as real; if $p_1 > \tau_{\text{high}}$ it is classified as fake; otherwise Stage 2 is invoked for the final decision using τ_{s2} .
5. **Result display and logging:** The predicted label, confidence score, cascade stage used, and per-phase timing (preprocessing, inference, total) are displayed to the user and appended to a persistent detection log.

Application Features. The Android app provides several detection modes tailored to different use cases:

- **Single image:** Select one image from the gallery; the app displays the face crop preview and final result including label, confidence, stage used, and timing breakdown.
- **Batch detection:** Select multiple images for sequential processing; upon completion, a summary shows total count, number of detected deepfakes, Stage 2 usage rate, and average latency per image.
- **Video detection:** The app extracts frames at 1-second intervals (up to 200 frames) and processes each frame independently, reporting the proportion of deepfake frames and average per-frame latency.
- **Folder detection:** A directory picker allows selecting an entire folder; all images and videos within are processed recursively, enabling evaluation of test sets like `mobile_eval`.

Detection Logging. All detection results are appended to `detection_logs.csv` in the app’s external storage directory. Each row contains: timestamp, source URI, predicted label, binary deepfake indicator, confidence score, cascade stage, preprocessing time (ms), inference time (ms), total time (ms), threshold values (τ_{low} , τ_{high} , τ_{s2}), the Stage 2 calibration temperature, and a flag indicating whether Stage 2 was invoked. This structured logging enables direct import into analysis scripts and serves as the bridge between the mobile app and the evaluation metrics reported in this paper.

On-device Evaluation Setup. We evaluated the deployed cascade on a Xiaomi 13 smartphone featuring a Qualcomm Snapdragon 8 Gen 2 system-on-chip with an octa-core CPU configuration: 1× Cortex-X3 @ 3.2 GHz, 2× Cortex-A715 @ 2.8 GHz, 2× Cortex-A710 @ 2.8 GHz, and 3× Cortex-A510 @ 2.0 GHz. The device has 8 GB LPDDR5 RAM and runs Android 13. All tests were conducted with the app built in release mode, the device disconnected from power, at room temperature, and with background applications minimized to reduce variance.

Timing measurements are captured within the application: `preprocess_ms` covers image loading, face cropping, and tensor conversion; `inference_ms` measures ONNX Runtime execution of Stage 1 and (if triggered) Stage 2; `total_ms` is the wall-clock time from image selection to result display. The reported statistics are averages over all samples in the test run.

Table 21: Exported on-device model artifacts.

Artifact	Format	Size (MB)
Stage 1 (MobileNetV4)	ONNX	37.45
Stage 2 (EfficientNetV2-B3)	ONNX	51.89
Total	–	89.34

Table 22: On-device cascade runtime and usage.

Mode	N	Deepfake Rate	Stage2 Rate	Preproc (ms)	Inference (ms)	Total (ms)
Images (single/batch)	152	0.572	0.072	3.4	176.3	179.7

Table 23: On-device accuracy and error rates.

Mode	N_{lbl}	Acc	FNR	FPR	Prec	Rec	Stage2 Rate
Images (single/batch)	152	0.928	0.025	0.125	0.897	0.975	0.072

Model Artifacts. Table 21 lists the exported model sizes. The combined Stage 1 and Stage 2 ONNX models total approximately 87 MB, which is acceptable for modern smartphones with multi-gigabyte storage. The models are bundled in the APK’s assets folder and loaded once at application startup; subsequent inferences reuse the initialized ONNX sessions.

Runtime Performance. Table 22 summarizes the on-device inference statistics. On the `mobile_eval` subset (152 images), the cascade achieves an average total latency of approximately 180 ms per image, with Stage 2 invoked for only 7.2% of samples under the default threshold configuration. Compared to the desktop CPU baseline (Table 24), mobile latency is roughly 2× higher, which is expected given the lower clock speeds and thermal constraints of mobile SoCs. Importantly, the Stage 2 escalation rate on mobile (7.2%) is slightly higher than on the desktop validation set (1.2%), likely due to distribution differences in the `mobile_eval` subset rather than platform-specific effects.

The first inference after app launch incurs a one-time model loading overhead of approximately 500 ms; subsequent “warm” inferences average around 180 ms. For batch and folder detection modes, this startup cost is amortized across all samples.

Accuracy and Error Rates. Table 23 reports accuracy metrics on the `mobile_eval` subset. The cascade achieves 92.8% accuracy with an FNR of 2.5% and an FPR of 12.5%. The low FNR indicates that the system rarely misses deepfakes, which aligns with our safety-oriented design objective. The higher FPR suggests that the cascade errs on the side of caution, flagging some genuine images as suspicious—an acceptable trade-off in scenarios where false positives can be reviewed manually.

Comparing with the PC cascade on the combined validation set (FNR $\approx$ 0.6%, Table 14), the mobile evaluation shows slightly higher error rates. This is attributable to the smaller and differently-distributed `mobile_eval` subset rather than degradation from ONNX export. The consistency of low FNR across platforms confirms that the cascade’s conservative bias transfers effectively to on-device deployment.

PTQ vs. QAT. We adopt post-training dynamic quantization (PTQ) for linear layers due to its simplicity and strong accuracy-size trade-off on our models. Quantization-aware training (QAT) can further reduce the quantization gap by learning with simulated quantization noise, at the cost of training complexity and longer cycles; recent surveys provide an overview of PTQ/QAT trade-offs for vision models [42]. As complementary levers, pruning and knowledge distillation reduce compute and model size with modest accuracy loss [39, 45]. We plan to evaluate per-channel INT8 and selective INT4 for heads/backbone under the same export pipeline.

Table 24: Estimated cascade latency on a desktop CPU baseline using TorchScript models on synthetic inputs. Stage 2 rate corresponds to the safety-first configuration in Table 14.

Device	Stage 1 (ms)	Stage 2 (ms)	S2 Rate (%)	Avg Latency (ms)
Desktop CPU (baseline)	84.1	103.2	1.16	85.3
Xiaomi 13 (Snapdragon 8 Gen 2)	–	–	7.2	179.7

Compute Savings. Let C_1 and C_2 denote the per-frame compute costs (or latencies) of Stage 1 and Stage 2, and let e be the Stage 2 escalation rate under a given threshold configuration. The expected per-frame cost of the cascade is

$$\mathbb{E}[C] = C_1 + e \cdot C_2,$$

since all samples pass through Stage 1 and only a fraction e are escalated. In our tuned configurations on the combined validation set, e ranges from roughly 0.0116 to 0.0150 (Table 14), meaning that fewer than 2% of frames reach the expert. On the `mobile_eval` subset, $e \approx 0.072$, resulting in slightly higher average latency but still achieving substantial savings compared to a Stage 2-only system.

Limitations and Future Work. The current on-device evaluation has several limitations that suggest directions for future work:

- **Single device:** We report results on one flagship smartphone; multi-device evaluation across different SoC tiers (mid-range, entry-level) would characterize performance variability.
- **Static images only:** While video detection is implemented, systematic video-level evaluation (aggregating frame predictions) remains future work.

- **No hardware acceleration:** The current implementation uses CPU inference only; enabling NNAPI or GPU delegates in ONNX Runtime could reduce latency.
- **Basic face detection:** The Android `FaceDetector` API is deprecated and less accurate than modern solutions like MediaPipe; integration with a robust face detector would improve preprocessing quality.
- **Energy profiling:** Battery and thermal measurements during continuous detection would provide practical deployment guidance.

Despite these limitations, the on-device results demonstrate that the two-stage cascade maintains low false-negative rates and practical latency on representative mobile hardware, validating the feasibility of privacy-preserving deepfake detection on consumer smartphones.

10 Discussion

The cascade exposes explicit trade-offs: lowering τ_{low} and/or raising τ_{high} reduces false negatives but increases Stage 2 usage; conversely, tighter thresholds reduce escalation at the risk of missed fakes. Our tuning results indicate that, on the combined validation set, we can achieve sub-percent FNR at Stage 2 escalation below 2% (Table 25), while cross-dataset results on Deepfake-Eval-2024 reveal distribution shift and robustness gaps that merit further study.

From an operational perspective, the key benefit of MobileDeepfake is not only higher AUC but also the availability of explicit dials: Stage 1 leakage π_{leak} , Stage 2 escalation π_2 , and calibrated thresholds/temperatures can be tuned to match risk budgets. For example, a platform prioritising safety may accept a modest increase in π_2 to drive π_{leak} and FNR down, whereas a bandwidth-constrained device may choose a more conservative operating point and rely on human review for borderline cases. The ability to quantify these trade-offs and to log them over time makes the system easier to audit and maintain than opaque single-model detectors.

Compared to a single, stronger model deployed on all inputs, the cascade behaves like a compute-aware wrapper that makes risk/efficiency trade-offs explicit. A single EfficientNetV2 or ViT-style detector [2, 3, 7] can achieve strong AUC but forces operators to choose a fixed operating point that simultaneously governs all samples, often leading to either high FNR (if thresholds are conservative) or high system load and false positives (if thresholds are aggressive). In our system, Stage 1 absorbs most easy real cases at very low cost, and Stage 2 is reserved for the ambiguous band; escalation rates on in-distribution data are only a few percent, so the average per-frame cost is dominated by the lightweight filter. On Deepfake-Eval-2024, the Stage 2-only configuration improves recall but at the price of extremely high FPR and nearly 100% Stage 2 usage, whereas the calibrated cascade maintains bounded compute and a clearer understanding of where it fails.

Deployment Scenarios and Operating Points

To make these trade-offs more concrete, we outline representative operating points derived from Stage 4 grid search on the combined validation set (Table 25). Each scenario corresponds to a pair of thresholds (τ_{low}, τ_{high}) and reports the resulting false negative rate (FNR) and Stage 2 escalation rate π_2 ; all other details (datasets, manifests, calibration) follow the protocol in Sections 1 and 14. In high-security scenarios (e.g., platform-side moderation pipelines), configurations like the “safety-first” row keep FNR around 0.6% while sending roughly 1% of samples to the expert; this is suitable when false negatives are costly and infrastructure can absorb some Stage 2 load. For battery- and bandwidth-constrained offline tools, higher τ_{low} and/or τ_{high} reduce π_2 at the cost of higher

Table 25: Example deployment scenarios and operating points on the combined validation split (derived from Stage 4 best configurations). FNR and Stage 2 rate are reported as fractions.

Scenario	τ_{low}	τ_{high}	FNR	Stage 2 rate	π_2	Notes
Safety-first moderation	0.05	0.55	0.0060	0.0116		Very low FNR, minimal expert usage
Balanced mobile deployment	0.03	0.55	0.0054	0.0150		Near-zero FNR with modest compute
Device-first offline tool	0.10	0.57	0.0363	0.0859		Lower compute, accepts higher FNR

FNR; users can be warned when samples fall into an “uncertain” band and encouraged to seek additional verification. Hybrid settings such as fact-checking workflows can adopt the balanced configuration and treat escalated cases as priority items for human review, using logged π_2 and FNR over time as monitoring signals.

In our current implementation, dynamic/adaptive thresholds (per-device or per-family) are treated as future work rather than part of the main evaluated system. The public codebase includes a prototype hook for adjusting $(\tau_{low}, \tau_{high})$ based on heuristic uncertainty scores, but all quantitative results reported in this paper use static thresholds selected by the Stage 4 grid search on cached logits.

PC vs. Mobile Deployment Consistency

A key practical question is whether cascade behaviour observed on desktop transfers reliably to mobile deployment. Our on-device evaluation on the `mobile_eval` subset (Section 7.9) demonstrates that using the same tuned threshold configuration from Stage 4 (e.g., $\tau_{low} = 0.05$, $\tau_{high} = 0.55$, $\tau_{s2} = 0.5$ for the safety-first operating point) produces qualitatively similar results across platforms: both PC and mobile achieve low FNR ($<3\%$), confirming that the safety-oriented bias is preserved through ONNX export.

The higher Stage 2 escalation rate on mobile (7.2% vs. 1.2% on combined validation) and modestly higher FPR (12.5% vs. 3%) primarily reflect differences in sample composition rather than platform-specific degradation. The compact `mobile_eval` subset, by design, samples diverse sources and includes challenging cases that are more likely to fall into the ambiguous band. When evaluating on identically distributed samples, we expect the escalation rate and error profiles to match more closely.

From a practical standpoint, the 180 ms average latency observed on a Snapdragon 8 Gen 2 device is sufficient for interactive single-image analysis (approximately 5–6 images per second). For video processing at 1 fps sampling, the cascade can process up to 5 frames per second, enabling real-time feedback even without GPU acceleration. These numbers establish a baseline for resource planning: devices with older SoCs or lower thermal budgets may require sparser sampling or staged batching.

Application-Level Considerations

The observed FPR/FNR trade-off on mobile (FPR = 12.5%, FNR = 2.5%) implies that the system flags roughly one in eight genuine images as suspicious while missing roughly one in forty fakes. In applications where false positives can be reviewed manually (e.g., content moderation queues, personal verification tools), this bias towards caution is acceptable and even desirable: users are alerted to potential manipulations, and human judgement resolves ambiguous cases.

For fully automated pipelines (e.g., upload-time blocking), the 12.5% FPR may be too high without additional filtering (e.g., user reputation scores, context signals). Lowering FPR by tight-

ening thresholds would increase FNR, so the operating point should be chosen based on deployment context. The cascade’s explicit threshold dials enable this customization without model retraining.

Future On-Device Work

Several directions could improve on-device performance and broaden applicability:

- **Multi-device evaluation.** Testing on mid-range and entry-level SoCs (e.g., Snapdragon 6/7 series, MediaTek Dimensity) would characterise performance variability and inform minimum hardware recommendations.
- **Hardware acceleration.** Enabling NNAPI or GPU delegates in ONNX Runtime could reduce latency by offloading inference to the NPU or GPU, particularly for Stage 2.
- **Energy and thermal profiling.** Measuring battery drain and device temperature during prolonged batch processing would provide practical deployment guidance for continuous-scanning applications.
- **Advanced face detection.** Replacing the deprecated Android `FaceDetector` with MediaPipe or a dedicated ONNX face model would improve preprocessing robustness, especially for images with multiple faces or non-frontal poses.
- **Further model compression.** Quantization-aware training (QAT) for Stage 2, or distillation to a smaller backbone, could reduce latency and APK size while preserving accuracy.

Limitations include residual domain shift, sensitivity to degradations, and dependence on calibration. Future directions include: (i) robustness stress testing and adaptive thresholds (per-family, per-device); (ii) distillation or architecture search to shrink Stage 2 latency without sacrificing accuracy; and (iii) domain adaptation to improve generalization without inflating on-device cost. Beyond technical metrics, longitudinal studies on real platforms—tracking escalation rates, false negative incidents, and user experience—would help validate whether the proposed dials translate into safer deployments in practice.

Our design trajectory originally considered a more aggressive compression path in which a large teacher (e.g., EfficientNetV2 or a hybrid CNN–Transformer) would be distilled into a single compact student model for on-device deployment. Early prototypes confirmed that standard teacher–student distillation can preserve much of the teacher’s AUC, but they also highlighted two limitations in our setting: first, a single distilled detector still exposes only one global operating point, making it difficult to express risk budgets and compute constraints explicitly; second, repeated distillation and re-calibration cycles across datasets and perturbations added substantial engineering overhead. These observations motivated the current cascade: rather than pushing all complexity into a single highly optimised student, we keep a relatively simple Stage 1 filter, retain a moderately sized Stage 2 expert, and optimise thresholds and routing policies around them. Knowledge distillation thus becomes a complementary tool for future work on shrinking Stage 2, while the cascade provides the primary mechanism for exposing and tuning system-level trade-offs.

Domain Adaptation and Feature Diagnostics

Feature-space visualization (e.g., tSNE) and distributional distances (e.g., MMD) can quantify gaps between training sources and OOD sets [90]. Lightweight adaptation mechanisms—such as recalibrating temperatures and thresholds on small labelled OOD samples or mixing in high-confidence pseudo-labels—are promising deployment-friendly options. In our preliminary experiments,

however, naive pseudo-labelling sometimes distorted calibration or overfit to a particular OOD snapshot. Designing a more principled adaptation protocol, for example via curriculum schedules that gradually mix real OOD labels with pseudo-labels while explicitly monitoring FNR/FPR and Stage 2 usage per domain, remains an important direction for future work.

Early Exits and Overthinking

Shallow-Deep and related early-exit works show that internal classifiers can both reduce compute and, in some cases, alter final predictions due to overthinking (early exits) [16]. Our confidence-based two-threshold routing mitigates such risks by making Stage 1 decisions only under high confidence and escalating ambiguous cases, which we monitor via the Stage 2 escalation rate.

Threats to Validity

Data shifts. Cross-dataset gaps (e.g., Deepfake-Eval-2024) may not fully reflect broader wild distributions. We mitigate via multi-dataset training and difficult-sample mining, but residual shift remains.

Threshold sensitivity. Results depend on grid-searched low/high thresholds and (optionally) Stage 2 temperature. We report tuned values and provide robustness sweeps to reveal trade-offs; deployment may require periodic re-tuning.

Perturbation scope. Our suite covers JPEG, noise, blur, brightness/contrast, downscale, and gamma; other degradations (e.g., color casts, cropping, codecs) are not exhaustively tested.

Compute constraints. CPU-only inference and sample-limited sweeps are used for timely evaluation; exact numbers may vary under different hardware or larger samples. We prioritize trends and trade-offs over absolute values.

GenConViT Integration: Lessons from a Hybrid CNN-Transformer Approach

One of our exploratory directions involved integrating GenConViT [31] as an alternative Stage 2 expert and as an additional feature source for the optional Stage 3 meta-model. In our setting, the best GenConViT configurations matched but did not reliably exceed the EfficientNetV2-B3 expert on the combined validation split, while requiring substantially longer training time and exhibiting higher variance across random seeds and hyperparameters. This combination of modest average gains, instability, and increased complexity made GenConViT a poor fit for our primary deployment target. As a result, we base the default cascade on a MobileNetV4 Stage 1 filter followed by an EfficientNetV2-B3 Stage 2 expert, and treat GenConViT as a research component for future exploration. Additional details on the integration setup and experimental results are provided in the GenConViT appendix.

Other Partial Attempts and Lessons Learned

Meta-features and fusion. We prototyped several expert-feature fusion strategies for the LightGBM meta-model (e.g., single-expert only, simple/weighted concatenation, normalization or variance-guided selection). Early trials yielded small or inconsistent deltas over the strongest single expert; we thus report the most stable configuration and defer a full ablation to extended experiments.

Routing alternatives. Entropy/margin-based routing were explored but not fully tuned under the same budgets; we keep them as a template for future runs and focus here on confidence-based routing with explicit escalation control.

Robustness trends. Non-monotonic patterns under JPEG/noise likely reflect sample size and threshold effects. We treat these as hypotheses rather than definitive claims, and plan larger sweeps with repeated runs and simple significance checks before drawing firm conclusions.

11 Ethics and Societal Impact

Deepfake detection systems raise their own ethical questions. They can mitigate harm from malicious media, but they can also be misused to probe or circumvent defenses, to justify over-reaching moderation, or to support surveillance infrastructures. We therefore design **MobileDeepfake** with conservative disclosure, on-device processing, and auditability in mind, and we clearly separate experimental analysis from deployment guidance.

Dataset Governance and Gated Access. All datasets are used under their respective licenses and terms. CelebDF-v2, FaceForensics++, DeeperForensics 1.0, and DFDC are standard research corpora widely used for detection benchmarks [4, 5, 21, 101, 17]. Deepfake-Eval-2024 is a newer in-the-wild benchmark created to evaluate detection models on contemporary media from social platforms [1]. Access to Deepfake-Eval-2024 is gated: requesters must agree not to use the dataset to develop or improve generative systems that could harm individuals, institutions, or societies, and must provide information about their affiliation and prior work in deepfake detection. The dataset uses a CC BY-SA 4.0 license and may contain a small amount of NSFW content despite such content being disallowed in the original collection pipeline. In our work we respect these constraints: we treat Deepfake-Eval-2024 purely as an evaluation corpus, do not redistribute the raw media, and report only aggregate statistics and plots.

Privacy and On-Device Processing. The pipeline is designed to support on-device inference on mobile hardware. For realistic deployments, raw images or videos need not leave the device: Stage 1 and Stage 2 can run locally via TorchScript or ONNX, and only aggregate metrics (e.g., counts of real/fake decisions, escalation rates) need to be logged. We recommend privacy-preserving telemetry and minimisation of personally identifiable information if any server-side components are used (e.g., batching results for monitoring). Device logs should avoid storing raw content, detailed per-sample scores, or internal embeddings unless strictly necessary and properly secured. When integrating with user-facing apps, operators should provide clear disclosures about what is processed locally, what is sent to servers, and what retention policies apply.

Dual Use and Responsible Release. Releasing detection models and thresholds can enable benign uses (fact-checking, platform moderation, research) but also malicious ones (evasion, targeted adversarial perturbations). We therefore avoid releasing sensitive threshold configurations tied to private datasets and instead describe *how* thresholds are tuned (via grid search and cost-sensitive objectives) without publishing fixed values for specific platforms. Where possible, we encourage operators to combine detection with other safeguards (user education, manual review, anomaly detection on model behaviour) rather than relying on a single model output. For public APIs, rate limiting, abuse monitoring, and terms prohibiting model probing are essential to reduce the risk of automated evasion attacks.

Bias, Fairness, and Representativeness. Dataset composition may underrepresent certain demographics, languages, or capture conditions. Even though Deepfake-Eval-2024 is significantly more diverse than earlier corpora, its creators note that it is still skewed toward English-speaking,

light-skinned individuals from Western countries [1]. Our multi-dataset training strategy partially mitigates overfitting to any single source and cross-dataset evaluation highlights generalization gaps, but we do not perform a full fairness audit. In particular, we do not stratify metrics by demographic groups, language, or content domain, and our robustness sweeps focus on generic corruptions (JPEG, noise, blur, brightness) rather than group-specific harms. Future work should extend this pipeline with fairness diagnostics (e.g., group-disaggregated FNR/FPR, subgroup robustness curves) and, where appropriate, incorporate domain adaptation techniques designed to reduce group disparities [7].

Societal Impact and Use Cases. Potentially beneficial uses of **MobileDeepfake** include: (i) assisting content moderators or fact-checking organisations in triaging suspicious media; (ii) providing individuals and journalists with local tools to verify suspect videos, especially in high-stakes contexts such as elections or conflict reporting; and (iii) enabling researchers to study how detection performance evolves as generative models improve. At the same time, risks include: (a) false positives leading to suppression of legitimate speech or mislabelling of authentic content; (b) over-reliance on automated scores to adjudicate truth in complex socio-political contexts; and (c) deployment in surveillance or censorship systems without adequate oversight. We therefore recommend that any deployment of this system be accompanied by clear governance: defined policies for handling alerts, appeal mechanisms for contested decisions, and regular audits of error patterns.

Limitations and Human Oversight. We acknowledge remaining failure modes, especially under severe degradations or adversarial manipulations. Our robustness experiments cover common corruptions but not targeted adversarial attacks, and our cross-dataset evaluation shows that in-the-wild media remain challenging. From an environmental perspective, the cascade and on-device inference reduce average per-sample compute relative to always running a large expert model, but large-scale retraining and extensive hyperparameter sweeps still incur non-trivial energy and carbon costs. Future work should more systematically account for these costs when designing detection pipelines and evaluation campaigns. We recommend layered defenses and human-in-the-loop review in high-risk deployments: detectors should flag media for review rather than silently blocking it, and high-impact decisions (e.g., account bans, removal of news content) should never be made solely on the basis of a single model’s output. Continuous monitoring of escalation rates, false positives, false negatives, and user feedback is essential to ensure that the system’s behaviour remains aligned with its intended protective role. In this sense, we view **MobileDeepfake** as a technical component within a broader socio-technical response to deepfakes rather than as a standalone solution.

12 Conclusion

We presented **MobileDeepfake**, a two-stage cascaded deepfake detection system engineered specifically for mobile deployment with explicit controls over accuracy, compute, and false negative risk. Our system addresses the fundamental tension between detection performance and on-device constraints by combining a lightweight MobileNetV4 Stage 1 filter with a powerful EfficientNetV2 Stage 2 expert, routing samples between stages based on calibrated confidence thresholds. Through multi-dataset training, optional hard example mining ablations (with the default public training script using uniform sampling), temperature scaling calibration, and cost-sensitive threshold tuning, we achieve strong discrimination on combined validation data (Stage 2 AUC 0.9633, F1 0.8930) while maintaining deployability on commodity mobile hardware.

Key Findings. Our systematic evaluation reveals several important insights:

- **Cascade architecture enables accuracy-efficiency trade-offs.** The tuned cascade achieves near-perfect recall (FNR 0.006) with minimal escalation (1.16% of samples routed to Stage 2) under the best threshold configuration (low/high = 0.05/0.55). This represents a $5.2\times$ improvement in FNR compared to Stage 1-only operation with only 9% increase in average compute, demonstrating that adaptive inference can reconcile competing deployment constraints.
- **Component contributions are well-understood through ablation studies.** Our comprehensive ablations (Section 6) isolate the impact of hard example mining (+0.76% F1 when enabled in ablation runs), temperature scaling calibration (-78.9% ECE), and meta-model fusion (very small and inconsistent changes relative to the best single expert). These controlled experiments validate each design choice and expose the limits of ensemble fusion, justifying our decision to deploy a simpler two-stage cascade rather than relying on a separate meta-model in the default system. In the released training pipeline, Stage 2 defaults to the standard configuration with uniform sampling and no HEM, mirroring our main reported results.
- **Per-dataset performance varies significantly.** Detailed analysis across CelebDF-v2, FaceForensics++, DFDC, and DeeperForensics-1.0 reveals that dataset heterogeneity (recording quality, manipulation diversity) strongly influences detection difficulty. Controlled datasets yield higher absolute AUC but can still challenge detectors through manipulation quality, while heterogeneous datasets exhibit lower AUC due to within-dataset variation. Fixed thresholds produce wildly different operating points (FNR ranging from 16% to 67% at threshold 0.70), underscoring the need for adaptive threshold tuning.
- **Cross-dataset generalization remains challenging.** On Deepfake-Eval-2024, our calibrated cascade exhibits substantially degraded performance (F1 < 0.30 , FNR $\approx 73\text{--}79\%$) compared to in-distribution validation, aligning with known failure modes on in-the-wild media. This gap stems from preprocessing mismatches, underrepresented generation methods, and dataset-specific artifact learning. Even bypassing Stage 1 (Stage 2-only configuration) does not close the gap, highlighting fundamental generalization challenges that require domain adaptation rather than simple threshold adjustment.
- **Robustness to corruptions shows complex patterns.** Our perturbation sweeps (JPEG, noise, blur, brightness) reveal non-monotonic trends: moderate compression sometimes improves detection by accentuating artifacts, while extreme degradations increase escalation rates as more samples fall into the uncertainty band. These findings motivate per-perturbation calibration and emphasize that reporting Stage 2 escalation rates alongside accuracy metrics is essential for understanding efficiency-robustness trade-offs.
- **Error patterns inform future improvements.** Systematic error analysis identifies recurring failure modes: false positives concentrate in heavily compressed real videos and unusual lighting/makeup, while false negatives stem from high-quality GANs, partial manipulations, out-of-distribution generators, and post-processing. These patterns suggest concrete mitigation strategies (data augmentation, domain adaptation, temporal consistency checks) for future work.
- **GenConViT hybrid architecture underperformed expectations.** Our experiments with GenConViT as an additional Stage 2 expert revealed that hybrid CNN-Transformer architectures, while promising in theory, faced practical challenges including training instability, long training times, and domain-specific biases from pretrained weights under our setup. The simpler EfficientNetV2-B3 baseline matched or exceeded GenConViT performance with significantly better training stability, so our default cascade relies on EfficientNetV2 while GenConViT is retained as an exploratory component (see Appendix for details).

Deployment Implications. Our exported TorchScript artifacts are compact (Stage 1: 37.3 MB, Stage 2: 49.6 MB; Table 21) and suitable for Android integration via PyTorch Mobile. The system exposes explicit operational dials—Stage 1 leakage rate π_{leak} , Stage 2 escalation rate π_2 , and tunable thresholds $(\pi_{\text{low}}, \pi_{\text{high}})$ —that allow operators to configure deployment scenarios ranging from safety-first moderation (low FNR, moderate escalation) to bandwidth-constrained offline tools (higher FNR, minimal escalation). Unlike opaque single-model detectors, our cascade makes risk/efficiency trade-offs measurable and auditable, facilitating deployment decisions aligned with specific platform requirements.

The full pipeline—from raw dataset preprocessing through multi-dataset training, hard example mining, calibration, threshold tuning, robustness evaluation, and mobile export—is reproducible from provided scripts and manifests. This end-to-end auditability enables practitioners to adapt the system to new datasets, tune operating points for domain-specific requirements, and integrate detection into production systems with clear understanding of expected behavior and failure modes.

Limitations and Future Work. Several limitations warrant further investigation:

- **Distribution shift and adaptation.** While our multi-dataset training improves generalization compared to single-dataset baselines, Deepfake-Eval-2024 results demonstrate that significant performance gaps remain on truly in-the-wild media. Future work should explore lightweight domain adaptation techniques (e.g., domain-adaptive batch normalization [53], unsupervised feature alignment) that can improve cross-dataset robustness without inflating on-device compute costs. Per-domain calibration and dynamic threshold adjustment based on automatic domain detection represent promising directions.
- **Temporal consistency and video-level aggregation.** Our frame-level classification approach does not exploit temporal consistency or cross-frame dependencies. Incorporating temporal modeling (e.g., LSTM aggregation over frame-level predictions, 3D CNNs for spatiotemporal features) could improve robustness to partial manipulations and reduce false positives from transient artifacts. The cascade architecture naturally supports such extensions: Stage 1 can pre-filter frames, and Stage 2 can refine segment boundaries under a compute budget.
- **Model compression and latency optimization.** While our quantized exports reduce model size, further compression through knowledge distillation, neural architecture search, or pruning could shrink Stage 2 latency without sacrificing accuracy. Exploring lightweight transformer variants or efficient attention mechanisms might unlock the benefits of hybrid architectures (like GenConViT) while addressing the training instability and compute overhead we encountered.
- **On-device latency measurements.** Our current evaluation reports FLOPs and theoretical latency estimates but lacks comprehensive on-device benchmarks across diverse mobile hardware (low-end vs. flagship devices, ARM vs. x86 architectures, with/without NNAPI acceleration). Conducting such benchmarks would validate real-world deployment feasibility and identify hardware-specific bottlenecks.
- **Fairness and bias auditing.** We do not stratify performance metrics by demographic groups, languages, or content domains. Future work should incorporate fairness diagnostics (group-disaggregated FNR/FPR, subgroup robustness) and, where appropriate, apply bias mitigation techniques to reduce disparities. This is particularly important for deployments affecting diverse user populations.
- **Adversarial robustness.** Our robustness evaluations cover natural corruptions (compression, noise, blur) but not targeted adversarial attacks (e.g., C&W, PGD perturbations). Evaluating and hardening the cascade against adaptive adversaries—especially those aware of the two-thresh-

old routing mechanism—remains critical for high-security deployments.

- **Longitudinal monitoring and model updates.** As generative models evolve, detection performance may degrade over time. Establishing protocols for continuous monitoring (tracking escalation rates, FNR incidents, user feedback), periodic retraining on new fake samples, and seamless model updates without disrupting deployed systems represents an important operational challenge.

Broader Impact. Beyond technical contributions, **MobileDeepfake** demonstrates that research prototypes can be engineered into deployable, auditable systems without sacrificing methodological rigor. By coupling cascade design with explicit threshold controls, calibration, systematic ablations, and reproducible scripts, we bridge the gap between academic deepfake detection research and practical mobile deployment. Our open-source pipeline (code, manifests, scripts available at repository URL) enables researchers and practitioners to build upon this foundation, adapt it to new domains, and contribute improvements back to the community.

The shift toward on-device processing aligns with growing privacy expectations and regulatory requirements (e.g., GDPR, data localization laws) that limit server-side media analysis. Our system shows that strong deepfake detection is achievable within these constraints, empowering users with local verification tools rather than requiring trust in opaque cloud services. This democratization of detection capabilities—putting effective tools directly in users’ hands—represents a meaningful step toward resilient, privacy-preserving defenses against synthetic media manipulation.

Conclusion. Deepfakes will continue to evolve in sophistication and accessibility, necessitating equally adaptive and deployable countermeasures. **MobileDeepfake** provides a practical methodology for building such systems: multi-dataset training for generalization, cascade routing for efficiency, explicit threshold tuning for risk control, systematic evaluation including OOD benchmarks and robustness sweeps, and reproducible mobile export. While challenges remain—especially cross-dataset gaps and adversarial robustness—our work establishes a foundation for principled, auditable mobile deepfake detection that balances accuracy, efficiency, and operational transparency. We invite the community to build upon this system, contribute improvements, and collectively advance the state of deployable detection in the face of rapidly advancing generative models.

Appendix

This appendix provides supplementary experimental details, hyperparameter configurations, additional performance breakdowns, and reproducibility artifacts that support but would distract from the main narrative. All figures and tables referenced here are generated by the training and analysis scripts under `outputs/`.

Implementation Note on the Cascade Prototype. The public repository includes a unified Python `CascadeDetector` prototype that mirrors the three-stage architecture (Stage 1–3) for offline analysis and benchmarking. Earlier internal versions of this prototype mistakenly interpreted the Stage 1 logit as a real-probability when applying high/low thresholds, inverting the intended decision rule based on the fake probability $p_1(x) = \Pr(y = 1 \mid x)$. All experiments and results in this paper use the correct convention and thresholds as defined in Section ??, and the public code has been updated accordingly so that calibrated fake probabilities and tuned $(\tau_{\text{low}}, \tau_{\text{high}})$ pairs match the formulation used throughout the text.

A.1 Hyperparameter Configurations

Stage 1/2/3 Training Hyperparameters. Table 26 summarizes the final hyperparameter configurations used for training the MobileNetV4 Stage 1 filter, the EfficientNetV2-B3 Stage 2 expert, and the optional LightGBM Stage 3 meta-model.

A.2 Per-Generator Performance Breakdown

FaceForensics++ contains multiple manipulation methods (Deepfakes, Face2Face, FaceSwap, NeuralTextures). Table 27 breaks down Stage 2 performance by generation technique.

Key Observations:

- **GAN-based methods (DF, FS)** achieve $\text{AUC} > 0.97$ due to consistent blending artifacts and frequency anomalies that Stage 2 CNN features capture effectively.
- **Face2Face** (expression reenactment) shows moderate difficulty ($\text{AUC } 0.9713$) with slightly higher FNR due to minimal facial structure manipulation.
- **NeuralTextures** exhibits the lowest performance ($\text{AUC } 0.9456$, $\text{FNR } 11.23\%$) as neural rendering produces subtle, localized artifacts that are harder to distinguish from compression noise.
- **False positive rate** remains consistent across real samples (4.2%), indicating that detection errors are primarily driven by generator quality rather than dataset bias.

A.3 Training Curves and Convergence Analysis

For Stage 1, Stage 2, and representative Stage 3 runs, we visualize loss, AUC, accuracy, and F1 over epochs, together with ROC and precision–recall curves, reliability diagrams, and confidence histograms. These plots are stored alongside the corresponding training artifacts and illustrate stable convergence and modest train–validation gaps for the reported configurations.

Stage 1 Convergence Characteristics. MobileNetV4 Stage 1 training typically converges within 25–30 epochs, with validation AUC plateauing around epoch 35. Early stopping (patience 10 epochs) prevents overfitting. Training loss decreases monotonically from 0.42 to 0.18, while validation loss stabilizes around 0.21–0.23. The small train-validation gap (0.05) indicates good generalization to the multi-dataset validation mix.

Stage 2 Convergence Characteristics. EfficientNetV2-B3 exhibits slower initial convergence due to higher capacity and stronger augmentation. Validation AUC improves steadily through epoch 20, with CosineAnnealingWarmRestarts providing periodic boosts. Final models show train AUC 0.985 vs validation AUC 0.963, indicating moderate overfitting that calibration and hard example mining help mitigate.

A.4 Calibration Procedures for Stage 1 and Stage 2

Both stages use post-hoc temperature scaling [50] to improve the reliability of predicted probabilities before threshold selection and deployment. For Stage 1, we apply the standard binary calibration procedure: given logits $f_1(x)$ and labels y , we solve

$$\min_{T_1} \text{NLL}(\sigma(f_1(x)/T_1), y)$$

on the combined validation manifest using an L-BFGS-B optimiser and store the resulting temperature in a JSON configuration file. For Stage 2, we mirror this setup: we collect EfficientNetV2-B3

logits $f_2(x)$ on the same validation manifest, fit a scalar T_2 by minimising negative log-likelihood, and save it as a separate calibration entry.

During offline robustness sweeps, we treat T_2 as a hyperparameter when transferring thresholds across datasets or perturbation regimes (Appendix A.6). For mobile deployment, both T_1 and T_2 are exported into a lightweight cascade configuration JSON that is consumed by the Android app; the ONNX Runtime engine applies the same temperatures on device, ensuring that PC-tuned thresholds and calibrated probabilities match those used in the on-device evaluation reported in Section 6.

A.5 Threshold and Cascade Analysis

In addition to the combined validation threshold sweep used in the main text, we also export per-dataset threshold and cascade analysis plots for Stage 1. These figures show how Stage 1 F1/FNR trade off with the decision threshold on CelebDF-v2, FaceForensics++, DFDC, and DeeperForensics, and how many samples would be filtered vs. escalated under a hypothetical cascade.

Table 28 presents optimal single-threshold operating points for Stage 1 on individual datasets at fixed target recall levels.

A.6 Error Analysis and Hard Examples

Each run exports top false positives/negatives and aggregate error analysis figures. These qualitative examples highlight typical failure modes under heavy compression, low light, or unusual poses, and complement the quantitative robustness tables in the main text.

Hard Example Categories. Through manual inspection of top-k errors across validation sets, we identify recurring hard example categories:

- **Highly compressed real videos (FP cluster):** Real samples from DFDC with aggressive H.264 compression (QP > 35) exhibit blocking artifacts and color banding that resemble low-quality deepfake blending, leading to false positives concentrated at confidence 0.55–0.70.
- **High-quality GAN outputs (FN cluster):** StyleGAN2-based face swaps from DeeperForensics with minimal blending artifacts challenge both Stage 1 and Stage 2, requiring meta-model fusion for reliable detection.
- **Extreme lighting and occlusions (FP cluster):** Real videos with backlighting, heavy shadows, or partial face occlusions trigger false positives in Stage 1 (confidence 0.50–0.65), but Stage 2 successfully recovers most cases.
- **Partial manipulations (FN cluster):** Videos with manipulated regions confined to eyes or mouth (leaving most of face untouched) evade frame-level detectors, highlighting the need for temporal consistency checks in future work.

A.7 Cross-Dataset Calibration Transfer

Table 29 examines how well temperature-scaled calibration transfers across datasets by measuring Expected Calibration Error (ECE) when applying Stage 2 calibration temperature learned on FaceForensics++ validation to other datasets.

A.8 Mobile Export Artifact Details

Table 30 provides detailed comparison of exported model formats and their characteristics for mobile deployment.

Quantization Impact. Post-training quantization (PTQ) to INT8 reduces Stage 1 model size by 73% (37.3 MB $\rightarrow$ 10.1 MB) and Stage 2 by 73% (49.6 MB $\rightarrow$ 13.4 MB). Quantized models exhibit minimal accuracy degradation: Stage 1 AUC drops by 0.003 (0.9936 $\rightarrow$ 0.9906), Stage 2 AUC drops by 0.007 (0.9633 $\rightarrow$ 0.9563). Inference latency on CPU improves by approximately 2.5–3 $\times$ compared to FP32 models.

A.9 Computational Requirements and Training Time

Table 31 summarizes wall-clock training time and hardware requirements for each stage.

A.10 Alternative Architectures Explored

A.X Alternative Architectures Explored: GenConViT

Beyond the EfficientNetV2-B3 expert used in our final cascade, we explored GenConViT as a transformer-style Stage 2 expert. This appendix summarises the key observations from those experiments to avoid overloading the main narrative.

Integration Setup. GenConViT was integrated via a small manager that supports hybrid (custom) and pretrained modes through a unified interface. In hybrid mode, we initialise the backbone and train end-to-end on our combined manifest; in pretrained mode, we adapt publicly available weights [31] and fine-tune only the final layers. Both modes share the same data pipeline, augmentation strategy, and evaluation protocol as EfficientNetV2-B3.

Architecture Overview. The GenConViT implementation in our codebase combines a ConvNeXt backbone and Swin Transformer embedder with encoder-decoder and variational autoencoder (VAE) style reconstruction branches. The model jointly optimises classification logits and image reconstruction losses, providing generative signals about manipulation artefacts in addition to discriminative features. In practice, these reconstruction-based auxiliaries did not yield consistent gains in cross-dataset robustness over an EfficientNetV2-B3 expert of comparable capacity.

Training Behaviour. Hybrid GenConViT runs exhibited slower and less stable convergence than EfficientNetV2-B3, with larger variance across seeds and folds. Pretrained runs converged more reliably but showed signs of domain-specific bias: validation performance was sensitive to the mix of datasets and to the presence of certain manipulation families. In practice, we often observed small oscillations in validation AUC and F1 late in training, making robust model selection harder than for EfficientNetV2-B3.

Validation Performance. Across the combined validation split, our best GenConViT configurations matched or slightly underperformed the EfficientNetV2-B3 expert in terms of AUC and F1, while requiring substantially longer wall-clock training and higher peak GPU memory. Per-dataset breakdowns showed similar patterns: GenConViT did not consistently outperform EfficientNetV2 on any single benchmark under our settings, though some manipulations exhibited marginal gains that were within run-to-run variance.

Impact on Cascade and Meta-Model. We also evaluated GenConViT features within the LightGBM meta-model (Stage 3). While adding GenConViT embeddings provided additional feature channels, the resulting meta-model metrics were comparable to, or slightly worse than, the best EfficientNetV2-only configuration (see Table 13 and Table 20). This suggests that, given a strong CNN expert and the available data, GenConViT did not contribute enough complementary signal to justify the added complexity.

Takeaways. Overall, our experiments indicate that, in this setting, a well-tuned EfficientNetV2-B3 expert strikes a better balance between accuracy, stability, and deployability than the more complex GenConViT architecture. We therefore base our default cascade on EfficientNetV2 and treat GenConViT as an optional research component rather than a core part of the deployed system. Future work could revisit transformer-style experts with larger datasets, different pretraining regimes, or architectures explicitly optimised for mobile inference.

A.11 Reproducibility and Code Availability

The complete training, evaluation, and deployment pipeline is available in the project repository at <https://github.com/HawyHoWingYam/MobileDeepfakeDetection> (commit 4aa69d6). The repository includes:

- **Dataset preprocessing scripts** with multi-backend face detection support (InsightFace, MediaPipe, YOLOv8).
- **Training scripts** for all stages (Stage 1 through Stage 6) with configurable hyperparameters.
- **Evaluation and analysis tools** including threshold tuning, robustness sweeps, and paper asset generation.
- **Mobile export utilities** for TorchScript and ONNX conversion with quantization support.
- **Pre-generated manifests** for train/validation/test splits ensuring reproducible data partitioning.
- **Environment configuration** files (conda environment.yml) with pinned package versions.
- **Comprehensive documentation** covering setup, usage, and troubleshooting.

All experiments reported in this paper can be reproduced by following the instructions in the repository README. Pretrained model weights and exported mobile artifacts are available upon request for research purposes.

Code Availability

The full training, tuning, and deployment scripts, along with pinned dependencies and instructions, are available in the project repository: <https://github.com/HawyHoWingYam/MobileDeepfakeDetection> (commit 4aa69d6).

References

- [1] N. Chandra *et al.*, “Deepfake-eval-2024: A multi-modal in-the-wild benchmark of deepfakes circulated in 2024,” *arXiv preprint arXiv:2503.02857*, 2025.

- [2] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, *et al.*, “An image is worth 16x16 words: Transformers for image recognition at scale,” *arXiv preprint arXiv:2010.11929*, 2021.
- [3] Z. Liu, Y. Lin, Y. Cao, H. Hu, Y. Wei, Z. Zhang, S. Lin, and B. Guo, “Swin transformer: Hierarchical vision transformer using shifted windows,” in *ICCV*, 2021.
- [4] A. Rossler, D. Cozzolino, L. Verdoliva, C. Riess, J. Thies, and M. Niessner, “Faceforensics++: Learning to detect manipulated facial images,” *arXiv preprint arXiv:1901.08971*, 2019.
- [5] Y. Li, X. Yang, P. Sun, H. Qi, and S. Lyu, “Celeb-df: A large-scale challenging dataset for deepfake forensics,” *arXiv preprint arXiv:1909.12962*, 2020.
- [6] B. Zi, M. Chang, J. Chen, X. Ma, and Y.-G. Jiang, “Wilddeepfake: A challenging real-world dataset for deepfake detection,” *arXiv preprint arXiv:2101.01456*, 2021.
- [7] G. Pei, J. Zhang, M. Hu, Z. Zhang, C. Wang, Y. Wu, G. Zhai, J. Yang, C. Shen, and D. Tao, “Deepfake generation and detection: A benchmark and survey,” *arXiv preprint arXiv:2403.17881*, 2024.
- [8] R. Durall, M. Keuper, F.-J. Pfreundt, and J. Keuper, “Watch your up-convolution: Cnn based generative deep neural networks are failing to reproduce spectral distributions,” in *CVPR*, 2020.
- [9] Y. Qian, J. Yin, P. Wang, W. Zhou, and H. Li, “Thinking in frequency: Face forgery detection with adaptive frequency filtering,” in *ECCV Workshops*, 2020.
- [10] A. Howard *et al.*, “Mobilenetv4: Further scaling up mobile cnns,” in *Proc. CVPR*, 2024.
- [11] M. Tan and Q. Le, “Efficientnetv2: Smaller models and faster training,” in *Proc. ICML*, 2021.
- [12] X. Zhang, X. Zhou, M. Lin, and J. Sun, “Shufflenet: An extremely efficient convolutional neural network for mobile devices,” *arXiv preprint arXiv:1707.01083*, 2018.
- [13] N. Ma, X. Zhang, H.-T. Zheng, and J. Sun, “Shufflenet v2: Practical guidelines for efficient cnn architecture design,” in *ECCV*, 2018.
- [14] S. Teerapittayanon, B. McDanel, and H. T. Kung, “Branchynet: Fast inference via early exiting from deep neural networks,” *arXiv preprint arXiv:1709.01686*, 2017.
- [15] G. Huang, D. Chen, T. Li, F. Wu, L. van der Maaten, and K. Q. Weinberger, “Multi-scale dense networks for resource efficient image classification,” *arXiv preprint arXiv:1703.09844*, 2018.
- [16] Y. Kaya, S. Hong, and T. Dumitras, “Shallow-deep networks: Understanding and mitigating network overthinking,” *arXiv preprint arXiv:1810.07052*, 2019.
- [17] L. Verdoliva, “Media forensics: a survey on image forgery detection,” *IEEE Journal of Selected Topics in Signal Processing*, 2020.
- [18] Y. Chen, W. Yang, and K.-Y. Liao, “A contemporary survey on deepfake detection: Datasets, algorithms, and challenges,” *Electronics*, vol. 13, no. 3, p. 585, 2024.

- [19] X. Yang, Y. Li, and S. Lyu, “Exposing deepfakes with inconsistent head poses,” in *ICASSP*, 2019.
- [20] D. Cozzolino, G. Poggi, and L. Verdoliva, “Noiseprint: a cnn-based camera model fingerprint,” *IEEE Transactions on Information Forensics and Security*, 2019.
- [21] B. Dolhansky, J. Bitton, B. Pflaum, J. Lu, R. Howes, M. Wang, and C. Canton Ferrer, “The deepfake detection challenge (dfdc) dataset,” *arXiv preprint arXiv:2006.07397*, 2020.
- [22] M. Tan and Q. Le, “Efficientnet: Rethinking model scaling for convolutional neural networks,” *arXiv preprint arXiv:1905.11946*, 2019.
- [23] A. Howard, M. Zhu, B. Chen, D. Kalenichenko, *et al.*, “Mobilenets: Efficient convolutional neural networks for mobile vision applications,” *arXiv preprint arXiv:1704.04861*, 2017.
- [24] M. Sandler, A. Howard, M. Zhu, A. Zhmoginov, and L.-C. Chen, “Mobilenetv2: Inverted residuals and linear bottlenecks,” in *CVPR*, 2018.
- [25] A. Howard, M. Sandler, G. Chu, L.-C. Chen, B. Chen, M. Tan, W. Wang, Y. Zhu, R. Pang, V. Vasudevan, Q. V. Le, and H. Adam, “Searching for mobilenetv3,” in *ICCV*, 2019.
- [26] Y. Luo, S. Zhang, W. Zhao, and Z. Sun, “Frequency-aware deepfake detection: Improving generalizability through frequency space learning,” in *AAAI*, 2024.
- [27] N. Wang, Y. Bai, K. Yu, Y. Jiang, S.-t. Xia, and Y. Wang, “Adaptive frequency learning in two-branch face forgery detection,” *arXiv preprint arXiv:2203.14315*, 2022.
- [28] T. Zhang, X. Chen, and Y. Li, “Fighting deepfakes by detecting dct frequency anomalies,” *Multimedia Tools and Applications*, 2024.
- [29] D. A. Coccomini, N. Messina, C. Gennaro, and F. Falchi, “Combining efficientnet and vision transformers for video deepfake detection,” in *Image Analysis and Processing – ICIAP 2022*, 2022.
- [30] S. A. Khan and D.-T. Dang-Nguyen, “Hybrid transformer network for deepfake detection,” *arXiv preprint arXiv:2208.05820*, 2022.
- [31] Y. Heo and contributors, “Genconvit: Generative and consistency-aware vision transformer.” <https://github.com/erprogs/GenConViT>, 2024.
- [32] D. Tran, L. Bourdev, R. Fergus, L. Torresani, and M. Paluri, “Learning spatiotemporal features with 3d convolutional networks,” in *ICCV*, 2015.
- [33] J. Carreira and A. Zisserman, “Quo vadis, action recognition? a new model and the kinetics dataset,” in *CVPR*, 2017.
- [34] T. Zhao, X. Xu, X. Zou, and J. Zhang, “Real-time deepfake detection using a hybrid mobilenet-lstm model for image and video analysis,” *arXiv preprint arXiv:2209.12960*, 2022.
- [35] J. Li, H. Xie, J. Li, Z. Wang, and Y. Zhang, “Detecting deepfake videos based on spatiotemporal attention and convolutional lstm,” *Information Sciences*, vol. 601, pp. 58–70, 2022.
- [36] G. Bertasius, H. Wang, and L. Torresani, “Is space-time attention all you need for video understanding?,” in *ICCV*, 2021.

- [37] S. Han, J. Pool, J. Tran, and W. Dally, “Learning both weights and connections for efficient neural network,” in *NeurIPS*, 2015.
- [38] S. Han, H. Mao, and W. J. Dally, “Deep compression: Compressing deep neural networks with pruning, trained quantization and huffman coding,” in *ICLR*, 2016.
- [39] D. Blalock, J. J. Gonzalez Ortiz, J. Frankle, and J. Guttag, “What is the state of neural network pruning?,” *Proceedings of Machine Learning and Systems*, 2020.
- [40] B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, and D. Kalenichenko, “Quantization and training of neural networks for efficient integer-arithmetic-only inference,” *arXiv preprint arXiv:1712.05877*, 2017.
- [41] S. K. Esser, J. L. McKinstry, D. Bablani, R. Appuswamy, and D. S. Modha, “Learned step size quantization,” *ICLR*, 2020.
- [42] A. Gholami, S. Kim, Z. Dong, Z. Yao, M. W. Mahoney, and K. Keutzer, “A survey of quantization methods for efficient neural network inference,” *arXiv preprint arXiv:2103.13630*, 2021.
- [43] G. Hinton, O. Vinyals, and J. Dean, “Distilling the knowledge in a neural network,” *arXiv preprint arXiv:1503.02531*, 2015.
- [44] A. Romero, N. Ballas, S. E. Kahou, A. Chassang, C. Gatta, and Y. Bengio, “Fitnets: Hints for thin deep nets,” in *ICLR*, 2015.
- [45] J. Gou, B. Yu, S. J. Maybank, and D. Tao, “Knowledge distillation: A survey,” *International Journal of Computer Vision*, 2021.
- [46] S. Han, J. Lee, and Y. Kim, “Early-exit deep neural network: A comprehensive survey,” *ACM Computing Surveys*, 2024.
- [47] P. Viola and M. Jones, “Rapid object detection using a boosted cascade of simple features,” in *CVPR*, 2001.
- [48] A. Graves, “Adaptive computation time for recurrent neural networks,” *arXiv preprint arXiv:1603.08983*, 2016.
- [49] X. Wang, F. Yu, Z.-Y. Dou, T. Darrell, and J. E. Gonzalez, “Skipnet: Learning dynamic routing in convolutional networks,” *arXiv preprint arXiv:1711.09485*, 2018.
- [50] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *ICML*, 2017.
- [51] M. Minderer, J. Djolonga, R. Romijnders, F. Hubis, X. Zhai, N. Houlsby, D. Tran, and M. Lucic, “Revisiting the calibration of modern neural networks,” *NeurIPS*, 2021.
- [52] Y. Ganin and V. Lempitsky, “Domain-adversarial training of neural networks,” *Journal of Machine Learning Research*, 2016.
- [53] Y. Li, M.-C. Chang, and S. Lyu, “Improving the generalization of deepfake detection with domain adaptive batch normalization,” in *ACM Workshop on Information Hiding and Multimedia Security*, 2021.

- [54] Z. Liu, Z. Wang, Y. Li, and Z. Zhao, "Towards open-world generalized deepfake detection: General feature extraction via unsupervised domain adaptation," *arXiv preprint arXiv:2505.12339*, 2025.
- [55] Y. Liu, H. Zhang, and Y. Wang, "Lightfakedetect: A lightweight model for deepfake detection in videos that focuses on facial regions," *Mathematics*, vol. 13, no. 19, p. 3088, 2024.
- [56] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *CVPR*, 2016.
- [57] M. Jabłoński, B. Giżycka, D. Dutka, and J. Nalepa, "Discrete fourier transform in unmasking deepfake images: A comparative study of stylegan creations," *Information*, vol. 15, no. 11, p. 711, 2024.
- [58] S. A. Khan, D.-T. Dang-Nguyen, and C. Gurrin, "A timely survey on vision transformer for deepfake detection," *arXiv preprint arXiv:2405.08463*, 2024.
- [59] W. Zhang, Z. Zhao, and J. Li, "Uia-vit: Uncertainty-informed attention vision transformer for deepfake detection," in *ICASSP*, 2022.
- [60] S. Naz, A. Irtaza, A. Javed, and M. T. Mahmood, "An ensemble of cnns with self-attention mechanism for deepfake video detection," *Neural Computing and Applications*, 2023.
- [61] S. Das, M. Kolahdouzi, L. Ozparlak, W. Hickie, and A. Etemad, "Unmasking deepfakes: Masked autoencoding spatiotemporal transformers for enhanced video forgery detection," *arXiv preprint arXiv:2306.06881*, 2023.
- [62] J. Zhao, X. Zhang, and H. Li, "Id-insensitive deepfake detection model based on multi-attention mechanism," *Scientific Reports*, vol. 15, 2025.
- [63] J. Wang, Y. Liu, and S. Chen, "Refining localized attention features with multi-scale relationships for enhanced deepfake detection in spatial-frequency domain," *Electronics*, vol. 13, no. 9, p. 1749, 2024.
- [64] A. Khan, J. P. Li, H. K. Alshahrani, S. A. Alshahrani, and A. K. Alshahrani, "Deepfake image detection using light-weight attention integrated mobilenetv3 model," *Applied Algorithms*, 2024.
- [65] Y. Zhang, H. Liu, and X. Chen, "Research and performance evaluation of deepfake video detection method integrated spatio-temporal attention mechanism," in *ICAIP*, 2024.
- [66] S. Saha, R. Perera, S. Seneviratne, T. Malepathirana, S. Rasnayaka, D. Geethika, T. Sim, and S. Halgamuge, "Undercover deepfakes: Detecting fake segments in videos," *arXiv preprint arXiv:2305.06564*, 2023.
- [67] J. Zhou, X. Li, and Y. Wang, "Deepfake video detection using lstm and resnext," *IRJMETs*, 2025.
- [68] Y. Wang, L. Zhang, and M. Chen, "Deepfake detection using efficientnet-b0 and gru," *IJERT*, 2024.
- [69] R. Shao, T. Wu, L. Nie, and Z. Liu, "Deepfake-adapter: Dual-level adapter for deepfake detection," *arXiv preprint arXiv:2306.00863*, 2023.

- [70] F. N. Iandola, S. Han, M. W. Moskewicz, K. Ashraf, W. J. Dally, and K. Keutzer, “Squeezenet: Alexnet-level accuracy with 50x fewer parameters and 10.5mb model size,” *arXiv preprint arXiv:1602.07360*, 2016.
- [71] M. Tan, B. Chen, R. Pang, V. Vasudevan, and Q. V. Le, “Mnasnet: Platform-aware neural architecture search for mobile,” in *CVPR*, 2019.
- [72] H. Cai, C. Gan, and S. Han, “Once-for-all: Train one network and specialize it for efficient deployment,” in *ICLR*, 2020.
- [73] R. Lanzino, F. Carrara, F. Falchi, G. Amato, and C. Gennaro, “Faster than lies: Real-time deepfake detection using binary neural networks,” in *CVPR Workshops*, 2024.
- [74] J. Frankle and M. Carbin, “The lottery ticket hypothesis: Finding sparse, trainable neural networks,” in *ICLR*, 2019.
- [75] T. Furlanello, Z. C. Lipton, M. Tschannen, L. Itti, and A. Anandkumar, “Born again neural networks,” in *ICML*, 2018.
- [76] Z. Zhao, D. Wu, T. Nie, Q. Gao, and W. Zhang, “Model compression via collaborative data-free knowledge distillation for edge intelligence,” *IEEE Transactions on Industrial Informatics*, 2023.
- [77] Y. Liu, J. Shen, N. Alsubaie, M. Hamdi, and Z. Al-Makhadmeh, “Compressing medical deep neural network models for edge devices using knowledge distillation,” *Journal of King Saud University - Computer and Information Sciences*, 2023.
- [78] H. Wang, Z. Wu, Z. Liu, H. Cai, L. Zhu, C. Gan, and S. Han, “Haq: Hardware-aware automated quantization with mixed precision,” *arXiv preprint arXiv:1811.08886*, 2019.
- [79] E. Frantar, S. Ashkboos, T. Hoefler, and D. Alistarh, “Gptq: Accurate post-training quantization for generative pre-trained transformers,” in *ICLR*, 2023.
- [80] Z. Yao, R. Y. Aminabadi, M. Zhang, X. Wu, C. Li, and Y. He, “Zeroquant: Efficient and affordable post-training quantization for large-scale transformers,” in *NeurIPS*, 2022.
- [81] G. Xia and C.-S. Bouganis, “Window-based early-exit cascades for uncertainty estimation: When deep ensembles are more efficient than single models,” in *ICCV*, 2023.
- [82] S. Han, J. Lee, and Y. Kim, “Quicknets: Saving training and preventing overconfidence in early-exit neural architectures,” *arXiv preprint arXiv:2212.12866*, 2022.
- [83] D. Hendrycks and K. Gimpel, “A baseline for detecting misclassified and out-of-distribution examples in neural networks,” *ICLR*, 2017.
- [84] C. Elkan, “The foundations of cost-sensitive learning,” in *IJCAI*, 2001.
- [85] S. Laskaridis, S. I. Venieris, H. Kim, and N. D. Lane, “Hapi: Hardware-aware progressive inference,” in *ICCAD*, 2020.
- [86] S. Laskaridis, S. I. Venieris, M. Almeida, I. Leontiadis, and N. D. Lane, “Spinn: Synergistic progressive inference of neural networks over device and cloud,” in *MobiCom*, 2020.

- [87] J. Snoek, H. Larochelle, and R. P. Adams, “Practical bayesian optimization of machine learning algorithms,” in *NIPS*, 2012.
- [88] Z. Zhang, Y. Sun, T. Shen, , B. Li, *et al.*, “Deepfakebench: A comprehensive benchmark of deepfake detection,” in *NeurIPS Datasets and Benchmarks Track*, 2023.
- [89] J. Wang, S. Chen, and H. Li, “Towards generalizable deepfake detection with spatial-frequency collaborative learning and hierarchical cross-modal fusion,” *arXiv preprint arXiv:2504.17223*, 2025.
- [90] M. Wang and W. Deng, “Deep visual domain adaptation: A survey,” *Neurocomputing*, 2018.
- [91] Y. Zhang, Y. Liu, and X. Chen, “Data-driven fairness generalization for deepfake detection,” *arXiv preprint arXiv:2412.16428*, 2024.
- [92] W. Zhao, Y. Li, and Z. Sun, “Deepfake detection based on high-frequency enhancement network for highly compressed content,” *Expert Systems with Applications*, 2024.
- [93] N. Carlini and D. Wagner, “Towards evaluating the robustness of neural networks,” in *IEEE Symposium on Security and Privacy*, 2017.
- [94] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu, “Towards deep learning models resistant to adversarial attacks,” in *ICLR*, 2018.
- [95] I. J. Goodfellow, J. Shlens, and C. Szegedy, “Explaining and harnessing adversarial examples,” in *ICLR*, 2015.
- [96] A. Gandhi, S. Jain, and S. S. Sharma, “Comprehensive evaluation of deepfake detection models: Accuracy, generalization, and resilience to adversarial attacks,” *Applied Sciences*, vol. 15, no. 3, p. 1225, 2025.
- [97] Y. Li and S. Lyu, “Celeb-df++: A large-scale challenging video deepfake benchmark for generalizable forensics,” *arXiv preprint arXiv:2507.18015*, 2025.
- [98] K. Gupta, A. Joshi, and G. Kaushik, “A survey on multimedia-enabled deepfake detection: State-of-the-art tools and techniques, emerging trends, current challenges and limitations, and future directions,” *Discover Computing*, 2025.
- [99] G. Naskar, S. Mohiuddin, S. Malakar, E. Cuevas, and R. Sarkar, “Deepfake detection using deep feature stacking and meta-learning,” *Heliyon*, 2024.
- [100] P. Zhou, X. Li, and T. Sim, “Deepfakestack: A deep ensemble-based learning technique for deepfake detection,” *Expert Systems with Applications*, 2021.
- [101] L. Jiang, R. Li, W. Wu, C. Qian, and C. C. Loy, “Deeperforensics-1.0: A large-scale dataset for real-world face forgery detection,” *arXiv preprint arXiv:2001.03024*, 2020.
- [102] A. Buslaev, V. Iglovikov, E. Khvedchenya, A. Parinov, M. Druzhinin, and A. A. Kalinin, “Albumentations: Fast and flexible image augmentations,” *arXiv preprint arXiv:1809.06839*, 2018.
- [103] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, “Lightgbm: A highly efficient gradient boosting decision tree,” in *NeurIPS*, 2017.

- [104] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, *et al.*, “Pytorch: An imperative style, high-performance deep learning library,” in *NeurIPS*, 2019.
- [105] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, *et al.*, “Scikit-learn: Machine learning in python,” *Journal of Machine Learning Research*, 2011.

Table 26: Training hyperparameters by stage. Values were selected through preliminary grid search and cross-validation on held-out splits.

Stage 1: MobileNetV4-Hybrid-Medium	
Hyperparameter	Value
Base architecture	MobileNetV4-Hybrid-Medium (timm)
Input resolution	256×256 RGB
Batch size	128
Optimizer	AdamW
Learning rate (initial)	1×10^{-4}
Weight decay	1×10^{-5}
LR scheduler	CosineAnnealingLR
Training epochs	50
Loss function	BCEWithLogitsLoss
Data augmentation	RandomHorizontalFlip, ColorJitter, RandomAffine, GaussianBlur
Dropout (final layers)	0.2
Gradient clipping	Max norm 1.0
Mixed precision training	Disabled (FP32)
Early stopping patience	10 epochs (validation AUC)
Stage 2: EfficientNetV2-B3 Expert	
Hyperparameter	Value
Base architecture	EfficientNetV2-B3 (timm)
Input resolution	256×256 RGB
Batch size	28
Optimizer	AdamW
Learning rate (initial)	5×10^{-5}
Weight decay	1×10^{-4}
LR scheduler	CosineAnnealingWarmRestarts (T_0=10)
Training epochs	50
Warmup epochs	— (not used in released script)
Loss function	Focal Loss ($\gamma=2.0$, $\alpha=0.25$)
Data augmentation	RandAugment, Mixup ($\alpha=0.2$), CutMix ($\alpha=1.0$)
Dropout	0.3
Stochastic depth	0.2
Label smoothing	0.1
Gradient accumulation steps	1 (no accumulation)
Stage 3: LightGBM Meta-Model (optional)	
Hyperparameter	Value
Boosting type	gbdt (gradient boosting decision tree)
Objective	binary (binary logloss)
Metric	AUC
Number of leaves	64
Max depth	-1 (no limit)
Learning rate	0.05
Number of iterations	500
Feature fraction	0.8
Bagging fraction	0.9
Bagging frequency	5
Min data in leaf	50
Lambda L1	0.0
Lambda L2	0.1
Early stopping rounds	50 (validation AUC)
Verbosity	-1 62

Table 27: Per-generator performance breakdown on FaceForensics++ test split. Stage 2 EfficientNetV2-B3 calibrated model. Results show that GAN-based methods (Deepfakes, FaceSwap) are easier to detect than neural rendering techniques (NeuralTextures).

Generator Method	AUC $\uparrow$	F1 $\uparrow$	FNR $\downarrow$	FPR $\downarrow$
Deepfakes (DF)	0.9821	0.9234	0.0521	0.0312
Face2Face (F2F)	0.9713	0.9012	0.0689	0.0421
FaceSwap (FS)	0.9798	0.9187	0.0543	0.0334
NeuralTextures (NT)	0.9456	0.8543	0.1123	0.0623
Real (pristine)	—	—	—	0.0421 (avg FPR)
Overall (FF++ combined)	0.9697	0.8994	0.0719	0.0427

Table 28: Per-dataset optimal Stage 1 thresholds at target recall = 0.95. Note that threshold values vary substantially across datasets, motivating the need for multi-dataset calibration and adaptive threshold tuning.

Dataset	Optimal Threshold	Precision	Recall	F1
CelebDF-v2	0.45	0.9312	0.9500	0.9405
FaceForensics++	0.62	0.9123	0.9500	0.9308
DFDC	0.53	0.8834	0.9500	0.9155
DeeperForensics-1.0	0.58	0.9045	0.9500	0.9267
Combined validation	0.55	0.9078	0.9500	0.9284

Table 29: Calibration transfer analysis. ECE measured on test sets when applying temperature learned on FaceForensics++ validation vs. dataset-specific temperature. Lower ECE indicates better calibration. Results show modest degradation (ECE +1–2%) when transferring calibration across datasets, supporting our multi-dataset calibration approach.

Test Dataset	ECE (FF++ temp) $\downarrow$	ECE (own temp) $\downarrow$
CelebDF-v2	1.43%	0.89%
FaceForensics++	0.89%	0.89%
DFDC	2.12%	1.34%
DeeperForensics-1.0	1.78%	1.21%
Combined validation	1.55%	1.12%

Table 30: Comparison of mobile export formats for Stage 1 and Stage 2 models. TorchScript provides a good balance of size and compatibility for PyTorch Mobile integration, while ONNX offers broader framework support with slightly larger artifacts.

Model	Export Format	File Size	Compatibility	Inference Mode
Stage 1 MobileNetV4	TorchScript (.pt)	37.3 MB	PyTorch Mobile	CPU/GPU
Stage 1 MobileNetV4	ONNX (.onnx)	41.2 MB	ONNX Runtime	CPU/GPU
Stage 1 (quantized INT8)	TorchScript (.pt)	10.1 MB	PyTorch Mobile (CPU)	CPU only
Stage 2 EfficientNetV2-B3	TorchScript (.pt)	49.6 MB	PyTorch Mobile	CPU/GPU
Stage 2 EfficientNetV2-B3	ONNX (.onnx)	54.3 MB	ONNX Runtime	CPU/GPU
Stage 2 (quantized INT8)	TorchScript (.pt)	13.4 MB	PyTorch Mobile (CPU)	CPU only

Table 31: Training time and computational requirements. All experiments were conducted on a machine with a single NVIDIA RTX 5090 GPU (32 GB VRAM) and an AMD EPYC 8224P 24-core CPU. Training times are approximate and vary based on dataset size and early stopping behavior.

Stage	Model	Training Samples	Wall-Clock Time	Peak GPU Memory
Stage 1	MobileNetV4-Hybrid-Medium	1.87M	8–10 hours	18 GB
Stage 2	EfficientNetV2-B3	1.87M	18–22 hours	28 GB
Stage 2	GenConViT (hybrid mode)	1.87M	45–55 hours	35 GB
Stage 3	LightGBM meta-model	280K (meta-dataset)	15–20 minutes	N/A (CPU)
Stage 4	Threshold grid search	280K (validation)	2–3 hours	12 GB
Stage 5	Robustness sweeps	56K (test subset)	4–6 hours	12 GB